

ETL Examples

-- Lab 01

-- Create a data base

```
CREATE DATABASE TIPS_SALES_PROD  
COMMENT = 'This is topper tips sales production database';
```

-- Create Schema

```
CREATE SCHEMA "TIPS_SALES_PROD"."SALES"  
COMMENT = 'This is TopperTips Sales Data';
```

-- Created a names internal stage

```
CREATE STAGE "TIPS_SALES_PROD"."SALES".SALES_STAGE  
COMMENT = 'This is named internal SF manage stage';
```

--Creating a stage via clone (this operation may not work with WebUI)

```
CREATE STAGE TIPS_SALES_PROD.SALES.S3  
CLONE USDA_NUTRIENT_STDREF.PUBLIC.S3_USDF_EXTERNAL_NAMED_STAGE  
COMMENT = 'Cloning USDF S3 Bucket';
```

-- Creating a part table

```
CREATE TABLE TIPS_SALES_PROD.SALES.PART  
CLONE SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.PART  
COMMENT = 'This is cloned table from TPC AdHoc SF1';
```

-- <Error> SQL compilation error: Cannot clone from a table that was imported from a share.

-- alternative is -- use create like

```
CREATE TABLE TIPS_SALES_PROD.SALES.PART LIKE SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.PART  
COMMENT = 'This is create like table from TPC AdHoc SF1';
```

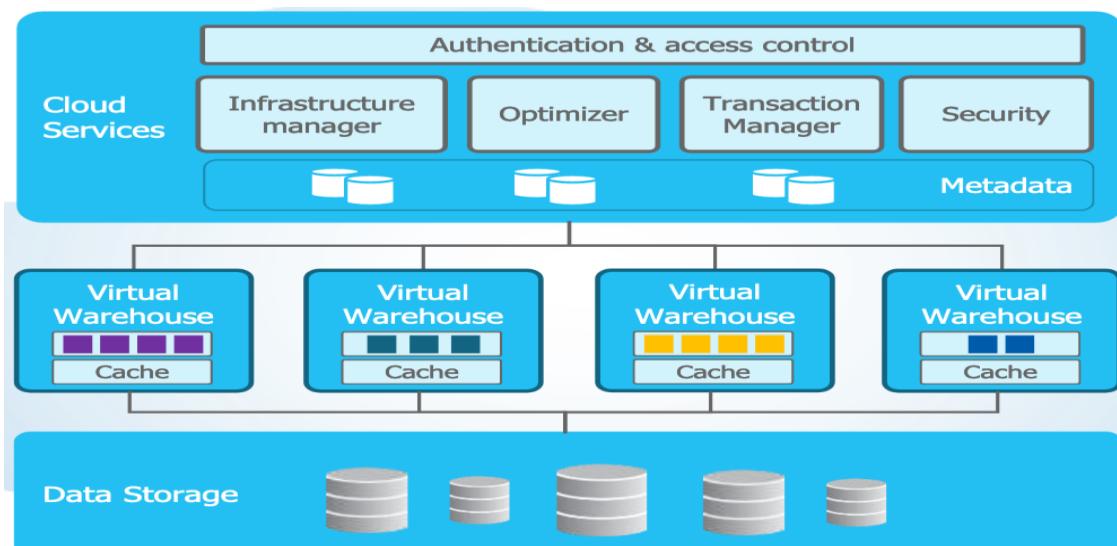
-- then insert

```
insert into TIPS_SALES_PROD.SALES.PART select * from  
SNOWFLAKE_SAMPLE_DATA.TPCH_SF1.PART
```

Three Layer

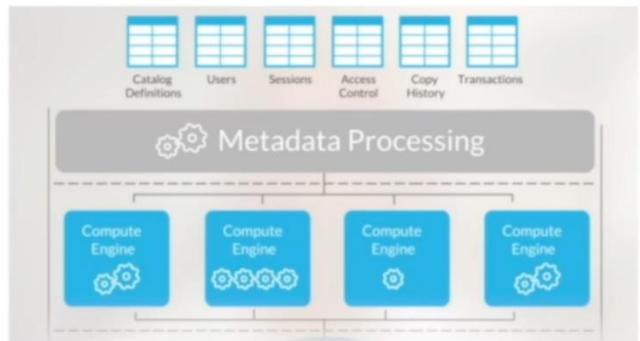
SNOWFLAKE ARCHITECTURE – CLOUD SERVICES

- Independent services managed entirely by Snowflake
Scalable, cost-effective cloud data warehouse architecture
- Built-in Fault Tolerance and High Availability (Continuous Availability)
Resilience against failures and delivers zero planned downtime
- Enterprise-grade Security
Complete industry security standards compliance and best practices
- Metadata-driven Cost-based Optimizer
Higher query performance through pruning and heuristic optimizations
- Time Travel and Zero-copy Cloning capabilities
Time-based data restoration and near instantaneous data snapshots
- Data Sharing
Secure, cross-account data access with absolutely no data movement
- Transaction Management with full ACID-compliance
Guaranteed consistent views of data with high concurrency and performance



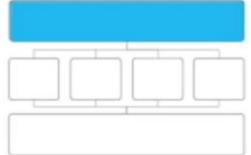
SNOWFLAKE METADATA SERVICE

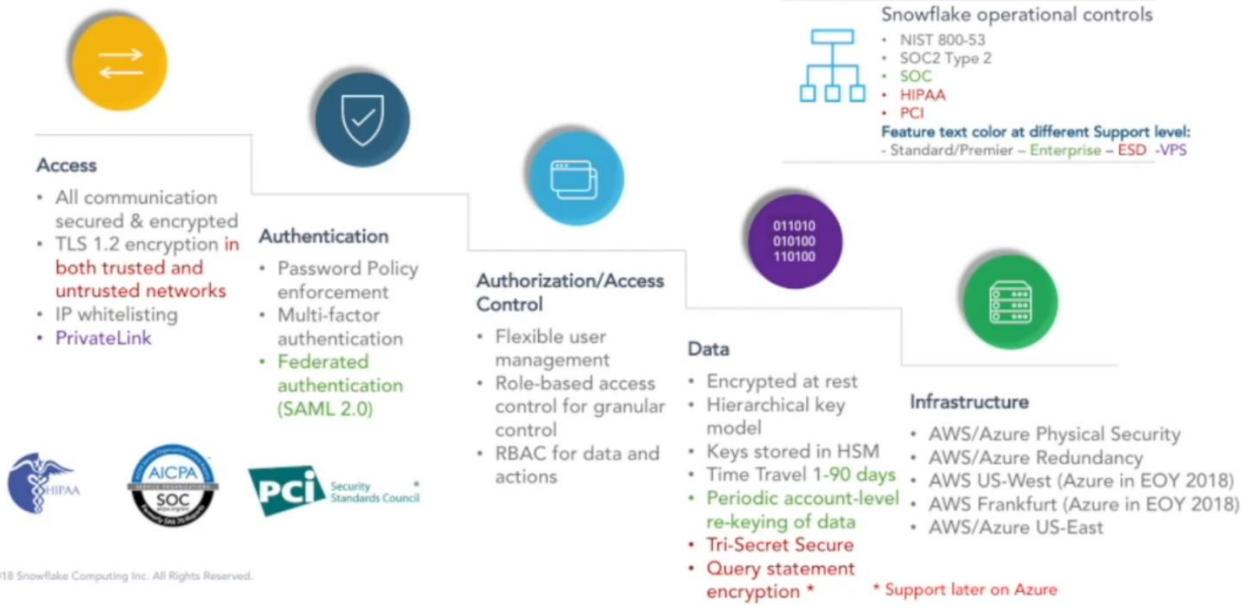
- Provides foundation for:
 - Optimized micro-partition pruning
 - Time travel
 - Zero-copy cloning
 - Data sharing
- Linear Scaling Performance
 - Metadata processing is separate from compute resources
- Every aspect of Snowflake architecture leverages metadata
 - All user visible actions have metadata backing
 - catalog definitions, accounts, users, sessions, databases, schemas, tables, roles, grants, ...
 - User invisible metadata such as:
 - Statistics for query optimizer
 - Encryption keys, servers, ...



SQL OPTIMIZER SERVICE

- SQL Optimizer [Cascades-style]
 - Top-down cost-based optimization (CBO)
- Automatic Statistics Gathering
 - On data loading and updates
- Built-in Join Order Optimization
 - No user input or tuning required
 - Use heuristics to automatically determine for more optimal join order for query performance
- Supports any and all data models
 - 3NF, snowflake, star, Data Vault
 - Optimizations for Star Schema (Bloom Filters, Broadcast Join)





Database Table View

1. No limit of db, db object in SF

2. default table is permanent table (temp, transient, external are other table types)

3. Non-permanent & transitory data can be stored in temporary & transient tables (short term data storage)

4. temporary & transient tables also called session specific data sets.

5. Temporary tables scope within the session in which they were created.

(so where the data is stored during the session)

- they don't support cloning
- not accessible via other sessions
- not recoverable (no fail safe) when session ends.
- SF also support temporary stages (other objects too)
- tmp data is part of storage cost and charges apply
- tmp table data can live more than 24hrs if session is alive and it cost storage charges.
- tmp table and permanent/transient table can have same name
- tmp table gets precedence over permanent & transient table
- this behaviour is given due to time travel (duplicate table name)
- tmp tables can not be converted to any other table once created
- retention period (time travel) is for 24 hours or the remainder of the session, whichever is shorter.

6. Transient Table

- There is a field in table (view) called is_transient
- Persist until explicitly dropped
- available to all users with privileges
- same as permanent table except fail safe
- contribute to the overall storage & no fail safe cost
- transient database and schemas are also possible
- schemas created under transient db and/or schema are by default transient.

- once created, can not be converted to other table type.
- it support clone

7. Compare

Type	Persistent	Cloning	TimeTravel	FailSafe(DR)
Permanent(Std)	Until Dropped	Yes	0 or 1	7
Permanent(Ent+)	Until Dropped	Yes	0 or 90	7
Transient	Until Dropped	Yes	0 or 1	0
Temporary	Session Level	Yes	0 or 1	0

8. timetravel can be specified for the table during creation or later via alter

9. The Fail-safe period is not configurable for any table type.

10. Transient and temporary tables have no Fail-safe period

11. External Table (ET)

- s3, azure, google cloud storage
- External tables store file-level metadata (file path, a version identifier, and partitioning information)
 - External tables support external (i.e. S3, Azure, or GCS) stages only; internal (i.e. Snowflake) stages are not supported.
 - Schema on read
 - No DML operation is allowed
 - External table can be used to perform join operation
 - Views can be created on external table
 - it is bit slower than the native SF tables
 - materialized views on external table improves performance (Enterprise+ edition)
 - Every ET has a column named VALUE of type VARIANT. Additional columns might be specified. All of the columns are treated as virtual columns.
 - * value (a variant type) - represent single row in external file
 - * METADATA\$FILENAME - a metadata column
 - METADATA\$FILENAME is a pseudocolumn, identifies stage data file and its path
 - virtual columns - it can be created using value and pseudocolumn

- No referential integrity constraints on external tables are enforced by Snowflake
- The default VALUE and METADATA\$FILENAME columns cannot be dropped.
- Not supported ****
 - * Clustering Key
 - * Cloning
 - * Data Sharing
 - * Data in XML format
 - * Time travel
- Recommendation for external table
 - * partition the table

12. Create External Table need following access

- Database - usage
- schema - usage/create stage/create external table
- stage - usage

-- Sample code to create a temp table and insert data .. all worked well

```
use database TIPS_SALES_PROD;
use schema public;
create temporary table t1 (f1 numeric);
desc table t1;
insert into t1 select p_size from part_expoert limit 100;
select * from t1;
drop table t1; -- this also works
```

-- if you run a show table, it shows the temporary tables too (refer image)

-- if you try to run following command on temporary table, you will get error

```
list @%t1
```

SQL compilation error: Stage 'TIPS_SALES_PROD.PUBLIC.%T1'" does not exist or not authorized.

```
drop table t2

create transient table t2 (parts_name varchar(1000)); --Table T2 successfully created.

insert into t2 select p_name from PART_EXPORT limit 1000; -- record inserted

select * from t2 limit 100;

show tables;

list @%t2; --this works
```

-- duplicate table name (refer images)

```
select * from PART_EXPORT limit 10; --this is a permanent table

create temporary table PART_EXPORT(parts_key number(38,0), parts_name varchar(55)); -- this
works

create transient table PART_EXPORT(parts_key number(38,0), parts_name varchar(55)); -- this does
not

drop table PART_EXPORT; -- this will drop the tmp table

show tables;
```

-- how to grant access to a database

```
GRANT USAGE, CREATE SCHEMA ON DATABASE "TIPS_DATABASE" TO ROLE "USERADMIN";
```

-- External Table Examples

-- additional key param is required

```
create or replace stage mystage url='s3://mybucket/files/'; -- aws

create or replace stage mystage url='gcs://mybucket/files'; -- gcp

create or replace stage mystage url='azure://myaccount.blob.core.windows.net/mycontainer/files'; -
- azure

-- blob.core.windows.net endpoint for all supported types of Azure blob storage accounts,
including Data Lake Storage Gen2.
```

-- create external table with stage reference

-- (Tip - create an internal stage reference and try to play with external table what happens)

-- desc external table

```
desc external table emp;
```

-- show external table

```
show external tables like 'line%' in tpch.public;
```

-- manually refresh external table

```
alter external table exttable_json refresh;
```

-- Similar to the first example, but manually refresh only a path of the metadata for an external table:

```
create or replace stage mystage
```

```
url='<cloud_platform>://twitter_feed/logs/'
```

```
..;
```

-- Create the external table

-- 'daily' path includes paths in </YYYY/MM/DD/> format

```
create or replace external table daily_tweets
```

```
with location = @twitter_feed/daily/;
```

-- Refresh the metadata for a single day of data files by date

```
alter external table exttable_part refresh '2018/08/05/';
```

-- Add an explicit list of files to the external table metadata:

```
alter external table exttable1 add files ('path1/sales4.json.gz', 'path1/sales5.json.gz');
```

-- Remove an explicit list of files from the external table metadata:

```
alter external table exttable1 remove files ('path1/sales4.json.gz', 'path1/sales5.json.gz');
```

```
-- drop external table
```

```
drop external table t2;
```

<https://s3.amazonaws.com/snowflake-docs/tutorials/json/server/2.6/2016/07/15/15/jsonTutorial.json>

```
{
  "device_type": "server",
  "events": [
    {
      "f": 83,
      "rv": "15219.64,783.63,48674.48,84679.52,27499.78,2178.83,0.42,74900.19",
      "t": 1437560931139,
      "v": {
        "ACHZ": 42869,
        "ACV": 709489,
        "DCA": 232,
        "DCV": 62287,
        "ENJR": 2599,
        "ERRS": 205,
        "MXEC": 487,
        "TMPI": 9
      },
      "vd": 54,
      "z": 1437644222811
    },
    {
      "f": 1000083,
      "rv": "8070.52,54470.71,85331.27,9.10,70825.85,65191.82,46564.53,29422.22",
      "t": 1437036965027,
      "v": {

```

```
    "ACHZ": 6953,  
    "ACV": 346795,  
    "DCA": 250,  
    "DCV": 46066,  
    "ENJR": 9033,  
    "ERRS": 615,  
    "MXEC": 0,  
    "TMPI": 112  
},  
    "vd": 626,  
    "z": 1437660796958  
}  
],  
"version": 2.6  
}
```

Data Type In Snowflake

Snowflake has its internal way of storing data, so on surface it support many style but many of them are just alias or same like int and decimal are all numbers

1. **Numeric** - Number (38,0) (rest all numbers are stored with this data type no matter what definition you give) = decimal,float,real, int, big int, etc...

2. String/Binary Data Type

= string/text/varchar
= char -> varchar(1)
= Binary = varbinary

3. Logical Data Type = Boolean (account supported post 25 Jan 2016)

4. Date/Time = Date

Time
DateTime
TimeStamp
TimeStamp_LTZ (Local Time Zone)
TimeStamp_NTZ (with no timezone)
TimeStamp_TZ (with timezone)

5. Semi Structured - Variant

Object
Array

Note : Snowflake displays FLOAT, FLOAT4, FLOAT8, REAL, DOUBLE, and DOUBLE PRECISION as FLOAT

Key Note about Numeric Data Type

- Numeric Data types are fixed point (int, bigint etc), floating point (float/double), constant/literal
- Number up to 38 with optional precision
- Precision - total number of digit allowed
- scale - totla # of digit right to the decimal point (37)
- data converted from higher to lower and back to higher, it will lose precision
- Example : convert a NUMBER(38,37) to DOUBLE (has 17 preision), and then back to NUMBER, it lose its precision
- Precision (total number of digits) does not impact storage <***IMP***>
- SF Micropartition determine min/max precision for given column and then store <***IMP***>
- Example: col values 1,2,3,4,5 -> 1 byte uncompressed
 - col values 1,2,3,100000 -> 4 byte uncompressed
- scale (number of digits following the decimal point) does have an impact on storage
- column of type NUMBER(10,5) consumes more space than NUMBER(5,0) + processing cost is higher

Key Note about Floating Point Numbers

- SF uses double-precision (64 bit) IEEE 754 floating point numbers.
- SF Supports following special values for Float
 - 'NaN' -> Not a number
 - 'inf' -> infinity
 - '-inf' -> negative infinity
 - ** these symbols must be single quote
- IEEE 745 vs SF when apply condition check
 - if('NaN' = 'NaN') .. in SF it is true but not in IEEE
 - IF('NaN' > X)in SF it is true but not in IEEE

Constant/Literals for Numbers

- e/E indicate exponent (At least one digit must follow the exponent marker if present.)
- Example 1.234E2 , 1.234E+2 , 15e-03
- +/- indicate positive or negative, default is +

Key Note about String & Binary Datatype

- VARCHAR holds unicode characters.
- max length is 16 MB (uncompressed) (single byte 16,777,216 and 2byte/4byte per char - 8,388,608 to 4,194,304)
- If a length is not specified, the default is the maximum length. <***IMP***>
- char/character same as varchar and not defined it is 1 default.
- for char datatype, no padding givein like standard database.

Key Note about Binary

- Max length = 8Mb
- No notion of unicode charater, so length is always same.
- varbinary = binary
- internal representation - The BINARY data type holds a sequence of 8-bit bytes.
- HELP is stored and will display as 48454C50 (H=48 in ascii)

String Constant

- String constants in SF must always be enclosed between delimiter characters.
- SF supports using either single quotes or dollar signs to delimit string constants <***IMP***>
 - Tip: in standard database, single quote is used for constant in select statement
- select 'today' as today, \$\$t o d a y\$\$ as today_with_dollar
- select \$\$wow \t wow\$\$ as today
- Two single quotes is not the same as the double quote character
- \$\$string with a ' character\$\$
- Note that the string constant cannot contain double-dollar signs. <***IMP***>

Boolean Data Type & Key Note

- BOOLEAN can have TRUE/FALSE values.
- BOOLEAN can also have an “unknown” value, which is represented by NULL
- Boolean Conversionn

*** String**

- * true/t/yes/y/on/1 = True

- * false/f/no/f/off/0 = False

* **Numeric**

- * 0 - zero is false

- * any non zero (+/-) is true

- String/Numeric conversionnnn

- * `to_string(True)` -> true (lower case) -> 1 (for number)

- * `to_string(False)` -> false (lower case) -> 0 (for number)

- Null is converted to NULL

Quik Review - Ternary Logic - three-valued logic (3VL)

- logic system with three truth values: TRUE, FALSE, and UNKNOWN

- Snowflake, UNKNOWN is represented by NULL

- When used in expressions (e.g. SELECT list), UNKNOWN results are returned as NULL values.

- When used as a predicate (e.g. WHERE clause), UNKNOWN results evaluate to FALSE.

Date/time key notes

- single DATE data type for storing dates (with no time elements)

- DATE supported formats for string constants (YYYY-MM-DD, DD-MON-YYYY, etc.)

- When a date-only format is used, the associated time is assumed to be midnight on that day.

- timestamp = combined date + time

- DATETIME is an alias for TIMESTAMP_NTZ

- all accepted timestamps are valid inputs for dates

- TIME data type for storing times in the form of HH:MI:SS

- TIME supports an optional precision parameter for fractional seconds, e.g. TIME(3)

- Time precision can range from 0 (seconds) to 9 (nanoseconds). The default precision is 9

- All TIME values must be between 00:00:00 and 23:59:59.999999999

- TIME internally stores “wallclock” time, and all operations on TIME values are performed without taking any time zone into consideration

calendar support

- Snowflake uses the Gregorian Calendar for all dates and timestamps

Date and Time Constants

- Snowflake supports using string constants to specify fixed date, time, or timestamp values
- String constants must always be enclosed between delimiter characters
- example

date '2010-09-14'

time '10:03:56'

timestamp '2009-09-15 10:59:43'

create table t1 (d1 date);

insert into t1 (d1) values (date '2011-10-29');

value format can be set as alter session set date-input-format ='yyyy-MM-dd'

Interval Constants

- interval constants to add/subtract a period of time to/from a date, time, or timestamp

- Example

INTERVAL '1 YEAR' represents 1 year.

INTERVAL '4 years, 5 months, 3 hours' represents 4 years, 5 months, and 3 hours

INTERVAL '1 year, 1 day' first adds/subtracts a year and then a day.

INTERVAL '1 day, 1 year' first adds/subtracts a day and then a year.

-SQL

```
select to_date ('2019-02-28') + INTERVAL '1 day, 1 year'; -- result 2020-03-01
```

```
select to_date ('2019-02-28') + INTERVAL '1 year, 1 day'; -- result 2020-02-29
```

- INTERVAL is not a data type and can be used only with date/time

-- More example

```
select to_date('2018-04-15') + INTERVAL '1 year'; -- 2019-04-15
```

```
select to_time('04:15:29') + INTERVAL '3 hours, 18 minutes'; --07:33:29
```

```
select name, hire_date from employees where hire_date > current_date - INTERVAL '2 y, 3 month';

select to_date('2025-01-17') + INTERVAL
      '1 y, 3 q, 4 mm, 5 w, 6 d, 7 h, 9 m, 8 s,
      1000 ms, 445343232 us, 898498273498 ns'
      as complex_interval2;

-- 2027-03-30 07:31:32.841
```

Simple Arithmetic for Dates

- TIME and TIMESTAMP values do not yet support simple arithmetic.
- select to_date('2018-04-15') + 1; -- 2018-04-16
- select to_date('2018-04-15') - 4; -- 2018-04-11

-- setting the timezone

```
alter session set timestamp_type_mapping = timestamp_ntz;
create or replace table ts_test(ts timestamp);
desc table ts_test;
```

```
create or replace table ts_test2(ts timestamp_ltz);
desc table ts_test2;
```

IFF Behavior

The **IFF** function returns the following results for ternary logic. Given a BOOLEAN column **C**:

If C is:	IFF(C, e1, e2)
evaluates to:	
TRUE	e1
FALSE	e2
NULL	e2

Logical Operators

Given a BOOLEAN column **C**:

If C is:	C AND NULL evaluates to:	C OR NULL evaluates to:	NOT C evaluates to:
TRUE	NULL	TRUE	FALSE
FALSE	FALSE	NULL	TRUE
NULL	NULL	NULL	NULL

In addition:

If C is:	C AND (NOT C) evaluates to:	C OR (NOT C) evaluates to:	NOT (C OR NULL) evaluates to:
TRUE	FALSE	TRUE	FALSE
FALSE	FALSE	TRUE	NULL
NULL	NULL	NULL	NULL

Run All Queries | Saving... 

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3           ' \ \t
4           \0x123_z $$,$$);
5 --below will fail
6 select $1, $2 from values ('row1', $$a
7           ' \ \t
8           \0x123 $$ $ $$);
```

Results Data Preview

Query ID SQL 16ms

SQL compilation error: syntax error line 3 at position 48 unexpected '\$'. parse error line 3 at position 54 near '<EOF>'.

Run All Queries Changes not saved ACCOUNTADMIN COMPUTE_WH

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3                                ' \ \t
4                                \0x123 z $ $$);
5
```

Results Data Preview

✓ Query ID SQL 55ms 1 rows Filter result... Copy

Row	\$1	\$2
1	row1	a' \ \t \0x123 z \$

Details

```
1 a
2                                ' \ \t
3                                \0x123 z $
```

Done

Run All Queries Changes not saved ACCOUNTADMIN COMPUTE_WH

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3                                ' \ \t
4                                \0x123 z $ $$);
5
```

Results Data Preview

✓ Query ID SQL 55ms 1 rows Filter result... Copy

Row	\$1	\$2
1	row1	a' \ \t \0x123 z \$

Run All Queries Saved 0 seconds ago ACCOUNTADMIN COMPUTE_WH

```
1 -- select from temporary table created using values statement
2 select $1, $2 from
3 values
4 ('Tab','He\tlo'),
5 ('Newline','He\nlo'),
6 ('Octal','-\041-'),
7 ('Hexadecimal','-\x21-')
8 ;
```

Results Data Preview

✓ Query ID SQL 362ms 4 rows Filter result... Copy

Row	\$1	\$2
1	Tab	He\tlo
2	Newline	He\nlo
3	Octal	-!-
4	Hexadecimal	-!-

Run All Queries | Saved 0 seconds ago

```
1 select 'today' as today
```

Results Data Preview

✓ Query ID SQL 51ms 1 rows

Filter result...

Row	TODAY
1	today

Run All Queries | Saving...

ACCOUNTADMIN COMPUTE_WH (XS) USDA_NUTRIENT... PUBLIC ...

```
1 -- create table with binary
2 create or replace table test_binary(
3     b100 binary(100),
4     vb varbinary
5 );
6 desc table test_binary;
7 --insert data
8 insert into test_binary
9 select to_binary('ab'), to_binary('AB'), to_binary('Ab')
10
11 select * from test_binary
12 -- insert into demo_binary (b) select to_binary(hex_encode('HELP')), 'HEX';
13
```

Results Data Preview Open History

✓ Query ID SQL 176ms 1 rows

Filter result... Columns ▾

Row	B	B100	VB
1	AB	AB	AB

Run All Queries | Saved 0 seconds ago

```
1 -- create table with binary
2 create or replace table test_binary(b binary,
3                                     b100 binary(100),
4                                     vb varbinary
5 );
6 desc table test_binary;
7 --insert data
8 insert into test_binary values ('HELP', 'help', 'Help')
9 |
10
```

Results Data Preview

✗ Query ID SQL 1.21s

String 'HELP' is not a legal hex-encoded string

```

1 -- create table with binary
2 create or replace table test_binary(b binary,
3                                     b100 binary(100),
4                                     vb varbinary
5                                     );
6 desc table test_binary;
7 --insert data
8 insert into test_binary values ('HELP', 'help', 'Help')
9
10

```

Results Data Preview

✓ Query ID SQL 47ms 3 rows

Filter result... [Download](#) [Copy](#)

Row	name	type	kind ↓	null?	default	primary key
1	B	BINARY(8388608)	COLUMN	Y	NULL	N
2	B100	BINARY(100)	COLUMN	Y	NULL	N
3	VB	BINARY(8388608)	COLUMN	Y	NULL	N

```

1 -- check the padding in character data type
2 create or replace table my_char_tab(
3 c char,
4 c1 char(1) ,
5 c10 char(10)
6 );
7 --insert values
8 insert into my_char_tab values ('a','b','c'),('x','y','zzzz')
9 --check length
10 select len(c), len(c1), len(c10) from my_char_tab

```

Results Data Preview

[Open History](#)

✓ Query ID SQL 201ms 2 rows

Filter result... [Download](#) [Copy](#)

Columns ▾ [X](#)

Row	LEN(C)	LEN(C1)	LEN(C10)
1	1	1	1
2	1	1	4

No padding done as it with other databases where char has padding for rest of the space

```

1 -- check the padding in character data type
2 create or replace table my_char_tab(
3 c char,
4 c1 char(1) ,
5 c10 char(10)
6 );
7 --insert values
8 insert into my_char_tab values ('a','b','c'),('xx','y','zzzz')

```

Results Data Preview

✗ Query ID SQL 518ms

String 'xx' is too long and would be truncated

```
1 -- constant example
2 create table my_constant(my_number number(38,30))
3 --insert data using constant/literals
4 insert into my_constant
5 values
6 (15),
7 (+1.34),
8 (0.2),
9 (15e-03),|
10 (1.234E2),
11 (1.234E+2),
12 (-1)
13 select * from my_constant
```

Important thing to notice here that 30 zeros appears after decimal due to definition

Results Data Preview

← Open History

✓ Query ID SQL 359ms 7 rows

 Copy

```
1 create or replace table test_boolean(
2   b boolean,
3   n number,
4   s string);
5 --insert some sample data
6 insert into test_boolean
7 values (true, 1, 'yes'),
8        (false, 0, 'no'),
9        (true, -1, 'on'),
10       (null, null, null);
11
12 -- casting nuber and string to boolean
13 select b col boolean, n col num, s col str, to boolean(n) numb to boolean, to boolean(s) str to boolean from test_boolean;
```

Results Data Preview

▶ Query ID SQL 1s 4 rows

 Copy

Row	COL_BOOLEAN	COL_NUM	COL_STR	NUMB_TO_BOOLEAN	STR_TO_BOOLEAN
1	TRUE		1 yes	TRUE	TRUE
2	FALSE		0 no	FALSE	FALSE
3	TRUE		-1 on	TRUE	TRUE
4	NULL		NULL NULL	NULL	NULL

ACCOUNTADMIN COMPUTE.WL (XS) USDA NUTRIENT PUBLIC

Results Data Previews

—Open Illia

✓ Query ID: SQL 146ms 3 rows

Row	name	type	kind	null?	default	primary key	unique key	check	expression	comment
1	B	BINARY(8388608)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
2	B100	BINARY(100)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
3	VB	BINARY(8388608)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL

```

1 create or replace table test_text(
2     v50 varchar(50),
3     c char,
4     c10 char(10),
5     s string,
6     s20 string(20),
7     t text,
8     t30 text(30)
9 );
10
11 desc table test_text;

```

Results Data Preview Open History

Query ID SQL 53ms 8 rows

Row	name	type	kind	null?	default	primary key	unique key	check	expression	comment
1	V	VARCHAR(16777216)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
2	V50	VARCHAR(50)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
3	C	VARCHAR(1)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
4	C10	VARCHAR(10)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
5	S	VARCHAR(16777216)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
6	S20	VARCHAR(20)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
7	T	VARCHAR(16777216)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
8	T30	VARCHAR(30)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL

▶ Run All Queries | Saved 3 seconds ago ACCOUNTADMIN COMPUTE_WH (XS) USDA_NUTRIENT... PUBLIC

```

1 create or replace table test_float(
2     f float,
3     dp double precision,
4     r real,
5     d_num number(38,8)
6 );
7
8 desc table test_float;
9 drop table test_float

```

Results Data Preview Open F

✓ Query ID SQL 59ms 5 rows

Row	name	type	kind	null?	default	primary key	unique key	check	expression	comment
1	D	FLOAT	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
2	F	FLOAT	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
3	DP	FLOAT	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
4	R	FLOAT	COLUMN	Y	NULL	N	N	NULL	NULL	NULL
5	D_NUM	NUMBER(38,8)	COLUMN	Y	NULL	N	N	NULL	NULL	NULL

Examples of Floating Point Data Types in a Table

```

create or replace table test_float(
    f float,
    dp double precision,
    r real
);

desc table test_float;

+-----+-----+-----+-----+-----+-----+-----+-----+
| name | type | kind | null? | default | primary key | unique key | check | expression | comment |
+-----+-----+-----+-----+-----+-----+-----+-----+
| D | FLOAT | COLUMN | Y | NULL | N | N | NULL | NULL | NULL |
| F | FLOAT | COLUMN | Y | NULL | N | N | NULL | NULL | NULL |
| DP | FLOAT | COLUMN | Y | NULL | N | N | NULL | NULL | NULL |
| R | FLOAT | COLUMN | Y | NULL | N | N | NULL | NULL | NULL |
+-----+-----+-----+-----+-----+-----+-----+-----+

```

Examples of Fixed-point Data Types in a Table

```
create or replace table test_fixed(num number,
                                   num10 number(10,1),
                                   dec decimal(20,2),
                                   numeric numeric(30,3),
                                   int int,
                                   integer integer
);

desc table test_fixed;

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| name | type      | kind   | null? | default | primary key | unique key | check | expression | comment |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| NUM  | NUMBER(38,0) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
| NUM10 | NUMBER(10,1) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
| DEC  | NUMBER(20,2) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
| NUMERIC | NUMBER(30,3) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
| INT  | NUMBER(38,0) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
| INTEGER | NUMBER(38,0) | COLUMN | Y     | NULL    | N          | N        | NULL  | NULL      | NULL    |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
```

Supported Date and Time Parts for Intervals

The INTERVAL keyword supports the following date and time parts as arguments (case-insensitive):

Date or Time Part	Abbreviations / Variations
year	y, yy, yyyy, yyyyy, yr, years, yrs
quarter	q, qtr, qtrs, quarters
month	mm, mon, mons, months
week	w, wk, weekofyear, woy, wy, weeks
day	d, dd, days, dayofmonth
hour	h, hh, hr, hours, hrs
minute	m, mi, min, minutes, mins
second	s, sec, seconds, secs
millisecond	ms, msec, milliseconds
microsecond	us, usec, microseconds
nanosecond	ns, nsec, nanosec, nsecond, nanoseconds, nanosecs, nseconds

```
6
7 create or replace table ts_test2(ts timestamp_ltz);
8
9 desc table ts_test2;
```

Results Data Preview

✓ Query ID SQL 57ms 1 rows

Filter result...

Row	name	type	kind	null?	default
1	TS	TIMESTAMP_LTZ(9)	COLUMN	Y	NULL

```
o  
7 create or replace table ts_test2(ts timestamp_ltz);  
8  
9 desc table ts_test2;
```

Results Data Preview

✓ Query ID SQL

57ms 57ms 1 rows

Filter result...



Copy

Row	name	type	kind	null?	default	primary key
1	TS	TIMESTAMP_LTZ	COLUMN	Y	NULL	N

▶ Run

All Queries | Saved 19 seconds ago

ACCOUNTADMIN

```
1 alter session set timestamp_type_mapping = timestamp_ntz;  
2  
3 create or replace table ts_test(ts timestamp);  
4  
5 desc table ts_test;
```

Results Data Preview

✓ Query ID SQL

68ms 68ms 1 rows

Filter result...



Copy

Row	name	type	kind	null?	default	primary key
1	TS	TIMESTAMP_NTZ(9)	COLUMN	Y	NULL	N

Variant Data Type

-- Variant/Array/Object Data Type

Semi-structure data type: used to represent arbitrary data structures

Can be used to import/operate on semi-structured data

JSON

Avro

ORC

Parquet

XML

Snowflake stores these types internally in an efficient compressed columnar binary representation

Variant

universal type, can store values of any other type, including OBJECT and ARRAY

maximum size of 16 MB compressed

A value of any data type can be implicitly cast to a VARIANT value, subject to size restrictions

example : (var:json_path >= 56) is cast to (var:json_path >= 56::VARIANT) for compare

VARIANT value can be missing (contain SQL NULL), which is different from a VARIANT null value

** VARIANT columns in a relational table are stored as separate physical columns

Object

Used to represent collections of key-value pairs,

key is a non-empty string, and the value is a value of VARIANT type

Snowflake does not currently support explicitly-typed objects.

Array

Used to represent dense or sparse arrays of arbitrary size

Unsupported Data Type

LOB (Large Object) - BINARY can be used instead

CLOB - VARCHAR can be used instead;

ENUM

User-defined data types

Querying Semi-structured Data

Snowflake supports SQL queries that access semi-structured data

JSON, Avro, ORC, and Parquet data (but not to xml)

Traverse Pattern Insert a colon : between the VARIANT column name and any first-level element

<column>:<level1_element>.

Query output is enclosed in double quotes because the query output is VARIANT, not VARCHAR

Operators : and subsequent . and [] always return VARIANT values containing strings

There are two ways to access elements in a JSON object:

Dot Notation

Bracket Notation

Regardless of which notation you use, the column name is case-insensitive but element names are case-sensitive

src:salesperson.name

SRC:salesperson.name

SRC:Salesperson.Name

Dot notation

```
select src:salesperson.name from car_sales;
```

Bracker notation -- Enclose element names in single quotes

```
select src['salesperson']['name'] from car_sales;
```

Repeating elements

```
select src:vehicle[0] from car_sales;  
select src:vehicle[0].price from car_sales;
```

Explicitly Casting

```
select src:salesperson.id::string from car_sales;
```

FLATTEN Function to Parse Arrays

FLATTEN is a table function that produces a lateral view of a VARIANT, OBJECT, or ARRAY

Example:

```
select  
    value:name::string as "Customer Name",  
    value:address::string as "Address"  
from  
    car_sales  
, lateral flatten(input => src:customer);
```

Json Data Load

#Sample JSON Data

```
[{"time":"6:38:51.000 PM","customer":"Talya McCambrois","action":"power off","device":"Footbot Air Quality Monitor"},  
 {"time":"8:13:16.000 PM","customer":"Willi Jenkerson","action":"power off","device":"GreenIQ Controller"},  
 {"time":"2:21:46.000 AM","customer":"Kippy Roux","action":"power off","device":"Amazon Echo"}]
```

```
CREATE DATABASE "TIPS_DATABASE"
```

```
CREATE FILE FORMAT "TIPS_DATABASE"."LAB".IOT_JSON  
TYPE = 'JSON'  
COMPRESSION = 'AUTO'  
ENABLE_OCTAL = FALSE  
ALLOW_DUPLICATE = TRUE  
STRIP_OUTER_ARRAY = FALSE  
STRIP_NULL_VALUES = TRUE  
IGNORE_UTF8_ERRORS = FALSE  
COMMENT = 'JSON File Format';
```

-- Step 4: Creating an internal stage to load data

```
CREATE STAGE "TIPS_DATABASE"."LAB".IOT_LAB_STAGE  
COMMENT = 'This is iot data loading stage';
```

-- Step 4: Creating a table

```
CREATE TABLE "TIPS_DATABASE"."LAB"."DEVICE_DATA" (  
"EVENT_TIME" TIMESTAMP NOT NULL,
```

```
"CUSTOMER_NAME" VARCHAR (100) NOT NULL,  
"DEVICE_ACTION" VARCHAR (50) NOT NULL,  
"DEVICE_NAME" VARCHAR (200) NOT NULL)  
COMMENT = 'This is customer iot device data';
```

PUT file://<file_path>/1.json @DEVICE_DATA/ui1589027218487

```
COPY INTO "TIPS_DATABASE"."LAB"."DEVICE_DATA"  
FROM @/ui1589027218487  
FILE_FORMAT = '"TIPS_DATABASE"."LAB"."IOT_JSON"'  
ON_ERROR = 'CONTINUE'  
PURGE = TRUE;
```

#since file has issue.. it threw following error

Unable to copy files into table.

**SQL compilation error: JSON/XML/AVRO file format can produce
one and only one column of type variant or object or array.**

Use CSV file format if you want to load more than one column.

```
CREATE TABLE "TIPS_DATABASE"."LAB"."DEVICE_DATA_VARIANT" ("DEVICE_EVENT" VARIANT NOT  
NULL) COMMENT = 'This is single column data';
```

-- Data loaded into single row

-- approach 2 to have in mutlipel rows

```
CREATE TABLE "TIPS_DATABASE"."LAB"."DEVICE_DATA_VARIANT_2" ("DEVICE_EVENT" VARIANT  
NOT NULL) COMMENT = 'Trying to load on individual rows';
```

```
CREATE FILE FORMAT "TIPS_DATABASE"."LAB".Device_Event_Format TYPE = 'JSON' COMPRESSION =  
'AUTO' ENABLE_OCTAL = FALSE ALLOW_DUPLICATE = TRUE STRIP_OUTER_ARRAY = TRUE  
STRIP_NULL_VALUES = TRUE IGNORE_UTF8_ERRORS = TRUE COMMENT = 'Each element as single  
row';
```

PUT file://<file_path>/1.json @DEVICE_DATA_VARIANT_2/ui1589027881962

```
COPY INTO "TIPS_DATABASE"."LAB"."DEVICE_DATA_VARIANT_2" FROM @/ui1589027881962  
FILE_FORMAT = "'TIPS_DATABASE"."LAB"."DEVICE_EVENT_FORMAT'" ON_ERROR = 'CONTINUE'  
PURGE = TRUE;
```

-- table to load ORC file via structured schema

```
CREATE TABLE "TIPS_DATABASE"."LAB"."ORC_DEVICE_DATA" ("EVENT_TIME" TIMESTAMP NOT  
NULL, "CUSTOMER_NAME" VARCHAR (100) NOT NULL, "EVENT_ACTION" VARCHAR (50) NOT NULL,  
"DEVICE_NAME" VARCHAR (200) NOT NULL) COMMENT = 'Storing Data from ORC File';
```

-- create file format to load data to table via put command

```
CREATE FILE FORMAT "TIPS_DATABASE"."LAB".ORC_DATA_FORMAT TYPE = 'ORC' COMMENT = 'This  
is ORC data format file';
```

-- data loading command via internal unmammed stage

**PUT file://<file_path>/part-00000-921b6fdc-41bb-4b06-b5d8-a8b760cb1558-c000.snappy.orc
@ORC_DEVICE_DATA/ui1589034165777**

```
COPY INTO "TIPS_DATABASE"."LAB"."ORC_DEVICE_DATA" FROM @/ui1589034165777  
FILE_FORMAT = "'TIPS_DATABASE"."LAB"."DEVICE_EVENT_FORMAT'" ON_ERROR =  
'ABORT_STATEMENT' PURGE = TRUE;
```

-- looks it also gave an error like other data set

Unable to copy files into table.

**SQL compilation error: JSON/XML/AVRO file format can produce one
and only one column of type variant or object or array.**

Use CSV file format if you want to load more than one column.

-- Now if you have lot of large data set then it will be an expensive operation

```
CREATE TABLE "TIPS_DATABASE"."LAB"."ORC_DEVICE_DATA_VARIANT" ("DEVICE_DATA" VARIANT  
NOT NULL) COMMENT = 'This table will have only one column';
```

```
PUT file://<file_path>/part-00000-921b6fdc-41bb-4b06-b5d8-a8b760cb1558-c000.snappy.orc  
@ORC_DEVICE_DATA_VARIANT/ui1589034420521
```

```
COPY INTO "TIPS_DATABASE"."LAB"."ORC_DEVICE_DATA_VARIANT" FROM @/ui1589034420521  
FILE_FORMAT = "'TIPS_DATABASE"."LAB"."ORC_DATA_FORMAT'" ON_ERROR = 'CONTINUE' PURGE =  
TRUE;
```

-- Now parquet type

```
CREATE FILE FORMAT "TIPS_DATABASE"."LAB".PARQUET_DATA_FORMAT TYPE = 'PARQUET'  
COMPRESSION = 'SNAPPY' BINARY_AS_TEXT = TRUE COMMENT = 'File Format for Parquet Data';
```

-- load command

```
PUT file://<file_path>/part-00000-d5a70322-eb9d-4e02-a85f-e678c1029748-c000.snappy.parquet  
@ORC_DEVICE_DATA/ui1589034857768
```

```
COPY INTO "TIPS_DATABASE"."LAB"."ORC_DEVICE_DATA"  
FROM @/ui1589034857768  
FILE_FORMAT = "'TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT'"  
ON_ERROR = 'CONTINUE'  
PURGE = TRUE;
```

it also ended in errorUnable to copy files into table.

**SQL compilation error: JSON/XML/AVRO file format can produce one and only one column of type
variant or object or array. Use CSV file format if you want to load more than one column.**

-- Now defineing a new table

```
CREATE TABLE "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_VARIANT" ("DEVICE_EVENT"  
VARIANT) COMMENT = 'The single column data loaded via Variant';
```

-- different options

```
PUT file://<file_path>/part-00000-d5a70322-eb9d-4e02-a85f-e678c1029748-c000.snappy.parquet  
@PARQUET_DEVIDE_DATA_VARIANT/ui1589035032039
```

--Do not load any data in the file

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_VARIANT" FROM  
@/ui1589035032039 FILE_FORMAT = "TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"  
ON_ERROR = 'SKIP_FILE' PURGE = TRUE;
```

--Stop loading, rollback and return the error

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_VARIANT" FROM  
@/ui1589035032039 FILE_FORMAT = "TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"  
ON_ERROR = 'ABORT_STATEMENT' PURGE = TRUE;
```

--Do not load any data in the file if the error count exceeds:

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_VARIANT" FROM  
@/ui1589035032039 FILE_FORMAT = "TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"  
ON_ERROR = 'SKIP_FILE_100' PURGE = TRUE;
```

--Continue loading valid data from the file

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_VARIANT" FROM  
@/ui1589035032039 FILE_FORMAT = "TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"  
ON_ERROR = 'CONTINUE' PURGE = TRUE;
```

-- Now trying to understand the object type

```
CREATE TABLE "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_OBJECT" ("EVENT_RECORD"  
OBJECT NOT NULL) COMMENT = 'This table has one field with object as field type';
```

-- loading data via put and copy command

```
PUT file://<file_path>/part-00000-d5a70322-eb9d-4e02-a85f-e678c1029748-c000.snappy.parquet  
@PARQUET_DEVIDE_DATA_OBJECT/ui1589035396167
```

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVIDE_DATA_OBJECT" FROM  
@/ui1589035396167 FILE_FORMAT = "TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"  
ON_ERROR = 'CONTINUE' PURGE = TRUE;
```

-- ended up with error -- not sure why

Unable to copy files into table.

SQL execution internal error: Processing aborted due to error 300010:3163851172; incident
9323353.

-- Now the final one is to use the array

```
CREATE TABLE "TIPS_DATABASE"."LAB"."PARQUET_DEVICE_DATA_ARRAY" ("DEVICE_EVENT" ARRAY NOT NULL) COMMENT = 'This is an experiment with Parquet with array';
```

```
PUT file://<file_path>/part-00000-d5a70322-eb9d-4e02-a85f-e678c1029748-c000.snappy.parquet @PARQUET_DEVICE_DATA_ARRAY/ui1589035563444
```

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVICE_DATA_ARRAY" FROM @/ui1589035563444 FILE_FORMAT = "'TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT'" ON_ERROR = 'CONTINUE' PURGE = TRUE;
```

-- end up with error -- could be due to the role issue as format is created by different role

Unable to copy files into table.

SQL execution internal error: Processing aborted due to error 300010:3163851172; incident 5140068.

-- Now following is given..lets see what happens

```
GRANT UPDATE, INSERT, DELETE, REBUILD, REFERENCES, SELECT, TRUNCATE  
ON TABLE "TIPS_DATABASE"."LAB"."PARQUET_DEVICE_DATA_ARRAY"  
TO ROLE "SYSADMIN" WITH GRANT OPTION;
```

-- Now role is changed

```
PUT file://<file_path>/tips.snappy.parquet @PARQUET_DEVICE_DATA_ARRAY/ui1589035944814
```

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_DEVICE_DATA_ARRAY" FROM @/ui1589035944814 FILE_FORMAT = "'TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT'" ON_ERROR = 'CONTINUE' PURGE = TRUE;
```

-- got the error like this

There was an error while trying to stage "tips.snappy.parquet".

Reason: SQL access control error: Insufficient privileges to operate on table stage 'PARQUET_DEVICE_DATA_ARRAY'

-- looks the internal stage has issues ??? looks a big issue -- looks the int stage has access issue

-- so create a table again with sysadmin

```
CREATE TABLE "TIPS_DATABASE"."LAB"."PARQUET_TYPE_ARRAY" ("DEVICE_EVENT" ARRAY NOT NULL) COMMENT = 'Parquet data loaded in array data type';
```

-- step-2

```
PUT file://<file_path>/tips.snappy.parquet @PARQUET_TYPE_ARRAY/ui1589036262859
```

```
COPY INTO "TIPS_DATABASE"."LAB"."PARQUET_TYPE_ARRAY" FROM @/ui1589036262859  
FILE_FORMAT = '"TIPS_DATABASE"."LAB"."PARQUET_DATA_FORMAT"' ON_ERROR = 'CONTINUE'  
PURGE = TRUE;
```

-- step-3 looks error again

Unable to copy files into table.

SQL execution internal error: Processing aborted due to error 300010:3163851172; incident 5508373.

-- @TOD - understand the object and array data type and why it is not working with parquet.

Data Unloading Export

-- Data Unloading & Export Function

We all know that Snowflake is cloud native warehouse system, data unloading and export function cost.

Hence it is very important to understand the data unloading practices and follow the guidelines.

There are various ways data can be exported effectively (like cloud resion, buckets, local box etc) and following topics will help us to explore them in detail.

Points to remember

1. Snowflake supports bulk export (i.e. unload) of data from a database table into flat, delimited text files
2. Use 'COPY INTO <location>' command (similar to load) the data from the Snowflake database table into one or more files in a Snowflake or external stage.

(Tip: it is always stages?)
3. Use the get command to download the file.

(Tip: so it is 2 stage process.. like a actual warehouse.. get the item at stage area and then get it in your truck)
4. Select query can also be used instead of 'COPY INTO <location>'

(Tip: look for code syntax)
5. Select query support full syntax including JSON style to download.
6. 'COPY INTO <location>' command provides option to specify single file or multiple (default SINGLE = FALSE)
7. Snowflake provides each file a unique name

(Tip: try this out to reinforce the understanding and file name patterns)
8. Snowflake prefixes the generated filenames with "data_".
9. Snowflake appends a suffix that ensures each file name is unique across parallel execution threads; e.g. data_stats_0_1_0.
10. MAX_FILE_SIZE copy option to specify each file

(Tip: mb or gb or what?)
11. Compression/Encryption is decided if stage is external or internal.
12. Location of stages - Internal (local), AWS s3, GCP bucket, Azure file

13. File formats - delited files (CSV/TSV), JSON, Parquet
 (Tip: not supported ORC, xml)
14. File encoding - UTF-8
15. Compression - gzip (default)/bzip2/Brotli/ZStandard
16. File Encryption - for internal stages - 128-bit (256-bit need configuration)
17. File Encryption - for external - user supplied key (possible only if it is provided)

Super Important Notes

1. An empty string is a string with zero length or no characters, whereas NULL values represent an absence of data
2. For CSV (delimited files)
 - 2.1 An empty string - ,,, to indicate that the string contains zero characters
 - 2.2 a NULL value is typically represented by two successive delimiters (e.g. ,,,)
 (Tip: get your hands on to reinforce theory and also try with TSV)
3. FIELD_OPTIONALLY_ENCLOSED_BY = 'character' | NONE - use this option to define behaviour (PREFERRED ONE****)
4. EMPTY_FIELD_AS_NULL = TRUE|FALSE option set to FALSE
5. NULL_IF = ('<string>',['<string>']) : When unloading data from tables: Snowflake converts SQL NULL values to the first value in the list.
 (Tip: see all the option in play .. refer images)
6. ESCAPE_UNENCLOSED_FIELD = \\N (default)

-- Setting the param for ui session with

-- worksheet to work with specific role/warehouse/db/schema

```
use role sysadmin;
use warehouse compute_wh;
use database TIPS_SALES_PROD;
use schema public;
```

```

-- populate data to a table from the example table

CREATE TABLE part_export LIKE snowflake_sample_data.tpch_sf1.part

INSERT INTO part_export select * from snowflake_sample_data.tpch_sf1.part


-- copying data to stage in csv format

copy into @export_stg/default_csv from part_export

-- all records loaded to stage and md5 check is also generated

-- the default_csv will be the prefix _0_0_0.csv.gz


-- Note: if you don't give a name and if stage is having data followign error will come

-- <error> Files already existing at the unload destination: @export_stg. Use overwrite option to
force unloading.

-- use remove function as alterntively

REMOVE @export_stg/export_stg/default_csv_0_0_0.csv.gz


-- Creating a table

create or replace table tips_export(
    id number(8) not null,
    first_name varchar(255) default null,
    last_name varchar(255) default null,
    city varchar(255),
    state varchar(255)
);

-- Populate the table with data

insert into tips_export (id,first_name,last_name,city,state)
values
(1,'Goutam','Shetty','Salt Lake City','UT'),
(2,'Rue','Shane','Birmingham','AL'),

```

```
(3,'John','Heck','Columbus','GA'),  
(3,'John','Heck','Columbus','GA'),  
(3,'Pick','Me','New York','NY'),  
(3,'Drop','You','Columbus','GA'),  
(3,'Hell','Nick','Columbus','GA'),  
(3,'Sim','Jack','Columbus','GA'),  
(3,'Jim','Raw','Columbus','GA'),  
(3,'Tium','Keo','New York','NY');
```

-- create a stage

```
CREATE STAGE export_stg COMMENT = 'my_simple_stage';
```

-- initiate a copy command and load data in json using selected query

```
copy into @export_stg  
from (select object_construct('id', id, 'first_name', first_name, 'last_name', last_name, 'city', city,  
'state', state) from tips_export)  
file_format = (type = json);
```

-- Data unload in parquet format

```
copy into @export_stg/myfile.parquet  
from (select id, first_name, last_name, city,state from tips_export)  
file_format=(type='parquet')  
header = true;
```

-- export multiple & compression

```
REMOVE @export_stg/file_size_0_0_0.csv.gz  
REMOVE @export_stg/file_size_0_0_1.csv.gz  
copy into @export_stg/file_size from part_export  
file_format = (type=csv compression='gzip')  
max_file_size=102400;
```

Run All Queries | Saving... A

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3           ' \ \t
4           \0x123 z $ $$);
5 --below will fail
6 select $1, $2 from values ('row1', $$a
7           ' \ \t
8           \0x123 $$ $ $$);
```

Results Data Preview

✖ Query ID SQL 16ms

SQL compilation error: syntax error line 3 at position 48 unexpected '\$'. parse error line 3 at position 54 near '<EOF>'.

Run All Queries | Changes not saved ACCOUNTADMIN COMPUTE_WH

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3           ' \ \t
4           \0x123 z $ $$);
5
```

Results Data Preview

✓ Query ID SQL 55ms 1 rows

Filter result...



Copy

Row	\$1
1	row1

\$2
a ' \ \t \0x123 z \$

Details

1 a	' \ \t
2	\0x123 z \$
3	

Done

Run All Queries | Changes not saved ACCOUNTADMIN COMPUTE_WH (X)

```
1 -- select from temporary table created using values statement with dollar quoted
2 select $1, $2 from values ('row1', $$a
3           ' \ \t
4           \0x123 z $ $$);
5
```

Results Data Preview

✓ Query ID SQL 55ms 1 rows

Filter result...



Copy

Row	\$1
1	row1

\$2
a ' \ \t \0x123 z \$

Run All Queries | Saved 0 seconds ago ACCOUNTADMIN COMPUTE_WH (XS) USDA_NUTRIENT... PUBLIC ...

```

1 -- select from temporary table created using values statement
2 select $1, $2 from
3 values
4 ('Tab','He\llo'),
5 ('Newline','He\nlo'),
6 ('Octal','-\041-'),
7 ('Hexadecimal','-\x21-')
8 ;

```

Results Data Preview

✓ Query ID SQL 362ms 4 rows

Filter result...

Row	\$1	\$2
1	Tab	He\llo
2	Newline	He\nlo
3	Octal	-!-
4	Hexadecimal	-!-

Run All Queries | Saving... ACCOUNTADMIN COMPUTE_WH (XS) USDA_NUTRIENT... PUBLIC ...

```

1 -- create table with binary
2 create or replace table test_binary(b binary,
3                                     b100 binary(100),
4                                     vb varbinary
5                                     );
6 desc table test_binary;
7 --insert data
8 insert into test_binary
9 select to_binary('ab'), to_binary('AB'), to_binary('Ab')
10
11 select * from test_binary
12 -- insert into demo_binary (b) select to_binary(hex_encode('HELP')), 'HEX';
13

```

Results Data Preview Open History

✓ Query ID SQL 176ms 1 rows

Filter result... Columns ▾

Row	B	B100	VB
1	AB	AB	AB

Run All Queries | Saved 0 seconds ago

```

1 -- create table with binary
2 create or replace table test_binary(b binary,
3                                     b100 binary(100),
4                                     vb varbinary
5                                     );
6 desc table test_binary;
7 --insert data
8 insert into test_binary values ('HELP', 'help', 'Help')
9
10

```

Results Data Preview

✗ Query ID SQL 1.21s

String 'HELP' is not a legal hex-encoded string

```

1 -- create table with binary
2 create or replace table test_binary(b binary,
3                                     b100 binary(100),
4                                     vb varbinary
5                                     );
6 desc table test_binary;
7 --insert data
8 insert into test_binary values ('HELP', 'help', 'Help')
9
10

```

Results Data Preview

✓ Query ID SQL 47ms 3 rows

Filter result... [Download](#) [Copy](#)

Row	name	type	kind ↓	null?	default	primary key
1	B	BINARY(8388608)	COLUMN	Y	NULL	N
2	B100	BINARY(100)	COLUMN	Y	NULL	N
3	VB	BINARY(8388608)	COLUMN	Y	NULL	N

```

32 copy into @export_stg/myfile.parquet
33 from (select id, first_name, last_name, city,state  from tips_export)
34 file_format=(type='parquet')
35 header = true;

```

Results Data Preview

✓ Query ID SQL 320ms 1 rows

Filter result... [Download](#) [Copy](#)

Row	rows_unloaded	input_bytes
1	10	1517

```

24
25 copy into @export_stg
26 from (select object_construct('id', id, 'first_name', first_name, 'last_name', last_name, 'city', city, 'state', state) from tips_export)
27 file_format = (type = json);
28
29 list @export_stg
30
31 Total Duration 108ms
32 Compilation 27ms
33 Execution 81ms
34 Compiling SQL 27ms
35 Executing (serial) 48ms
36 Receiving request from client 33ms

```

✓ Query ID SQL 108ms 1 rows

Row	name	size	md5	last_modified
1	export_stg/data_0_0.json.gz	224	01e1c7f53ff04157d75de142a4c94901	Thu, 14 May 2020 08:05:09 GMT

```

25
26 copy into @export_stg
27 from (select object_construct('id', id, 'first_name', first_name, 'last_name', last_name, 'city', city, 'state', state) from tips_export)
28 file_format = (type = json);
29 list @export_stg

```

Results Data Preview

✓ Query ID SQL 108ms 1 rows

Filter result... [Download](#) [Copy](#) Columns ▾

Row	name	size	md5	last_modified
1	export_stg/data_0_0.json.gz	224	01e1c7f53ff04157d75de142a4c94901	Thu, 14 May 2020 08:05:09 GMT

Snowflake Security Overview

VPC - Virtual Private Cloud (Virtual Private Snowflake VPS)

Physical Security

24hr arm guards in data center

video surveillance data center

no access to any unauthorized personality to data center

not Snowflake personal nor Snowflake customer have access to these data centers

Data Redundancy

Provided by cloud provider

Network Access

Network Policy helps to control the SF access

IP Whitelisting - policy can be created to allow or disallow IPs

private link - private tunnel between customer and cloud provider (ESD or VPS customer)

Account Access & Authentication

MFA - Multi Factor Authentication can be implemented to increase Security

MFA is provided by Duo Security services

Once MFA is installed, Duo app to be installed by user

Each user must enable MFA by themselves.

All users with account admin role should have MFA enabled.

SSO (SAML 2.0) allows user to access via federated services (IDP Identity provider)

As long as IPD session is active, they can access SF

SSO/IDP is available enterprise edition & +

Object Security

All the objects (warehouse/db/schema/table etc) can be controlled by DAC/RBAC

DAC- Discretionary Access control

RBAC - Role Based Access control

SF implements hybrid model of DAC & RBAC

DAC handle the ownership, each object has an owner and owner has full access to the object.

RBAC - Handle all other access except ownership like object privilege and role access

Object privilege assign to role which are intern assign to users.

Data Security

All data is encrypted uses AES-256 strong Encryption

All files stored in stage area is automatically Encryption using AES-128 or AES-256

Special edition of SF allows periodic re-key and customer manage encryption.

Connectivity Security

All communication over internet is via HTTPS.

All communication is secure and encrypted via TLS 1.2 or higher

Compliance Security

Third Party

HIPPA

PCI (Payment Card for Industry for data security)

NIST 800-53

SOC

SOC-2 type II

SIG Lite (SIG Assessment) (Standardized Information Gathering - SIG Questionnaire Tools allow organizations to build, customize, analyze and store vendor questionnaires)

Application Activity Log

History tab provide all historical commands

All details including session id etc can be viewed.

Each query has query id which helps for trouble shooting (no access to data to even SF users)

User Access Audit Log

Login_History family of table function can be used to query login attempts

Query History

Query History is available for 7 days, it can be stored in SF table or external system.

Infrastructure Monitoring

SF users Threat Stack & Sumo Logic to monitor production Infrastructure

Lacework for behavioural Monitoring (use activity/network traffic/binaries)

All alerts are viewed by SF security team.

Penetration Testing

SF performs 7-10 penetration testing per year

Application Penetration Test

Network Penetration Test

Functional Penetration Test

All logs and findings are tracked to closure

Test results are available with customer under NDA

Data Cloning

-- Cloning a database

```
CREATE DATABASE TIPS_DATABASE_QA  
CLONE "TIPS_DATABASE"  
COMMENT = 'This is cloned database from Tips Database';
```

-- Create like feature to copy table, but data is not copied

```
CREATE TABLE "TIPS_DATABASE_QA"."LAB".DEVICE_DATA_CREATE_LIKE  
LIKE "TIPS_DATABASE_QA"."LAB"."DEVICE_DATA"  
COMMENT = 'Create like feature is used which is similar to clone';
```

-- Cloning a file format

-- img exist

```
CREATE FILE FORMAT "TIPS_DATABASE_QA"."LAB".CLONED_DEVICE_EVENT_FORMAT  
CLONE "TIPS_DATABASE_QA"."LAB"."DEVICE_EVENT_FORMAT"  
COMMENT = 'Cloning a file formate to the same schema';
```

-- Create a stage via clone.

-- img exist

```
CREATE STAGE "TIPS_DATABASE_QA"."LAB".CLONE_JSON_STAGE  
CLONE "TIPS_DATABASE_QA"."LAB"."JSON_STAGE"  
COMMENT = 'This is stage cloning for named external stage';
```

Time Owner Comment

Clone File Format

Name * CLONED_DEVICE_EVENT_FORMAT

Source LAB.DEVICE_EVENT_FORMAT

Comment Cloning a file formate to the same schema

Show SQL Cancel Finish

Results Data Preview [Open History](#)

View: SNOWFLAKE.ACCOUNT_USAGE.DATABASES [Data](#) [SQL](#)

Cloned Database has not yet appeared. it will appear when additional storage operation will be performed

Row	DATABASE_ID	DATABASE_NAME	DATABASE_OWNER	IS_TRANSIENT	COMMENT	CREATED	LAST_ALTERED	DELETED	RETENTION_TIME
1	6	USDA_NUTRIENT_STDREF	SYSADMIN	NO	USDA Nutrie...	2020-05-02 ...	2020-05-02 ...		1
2	3	DEMO_DB	SYSADMIN	NO	demo databa...	2020-05-02 ...	2020-05-02 ...		1
3	4	SNOWFLAKE_SAMPLE_DATA	ACCOUNTAD...	NO	TPC-H, Open...	2020-05-02 ...	2020-05-02 ...		1
4	8	DB_FOR_SECURITY	SYSADMIN	NO	This is testing	2020-05-11 2...	2020-05-11 2...		1
5	2	UTIL_DB	SYSADMIN	NO	utility databa...	2020-05-02 ...	2020-05-02 ...		1
6	5	MAY022020DB	SYSADMIN	NO	This is sampl...	2020-05-02 ...	2020-05-02 ...		1
7	1	SNOWFLAKE		NO		2020-05-02 ...	2020-05-02 ...		1
8	7	TIPS_DATABASE	SYSADMIN	NO	This is practi...	2020-05-09 ...	2020-05-09 ...		1

Results Data Preview [Open History](#)

View: TIPS_DATABASE_QA.INFORMATION_SCHEMA.TABLES

Row	TABLE_CATALOG	TABLE_SCHEMA	TABLE_NAME	TABLE_OWNER	TABLE_TYPE	IS_TRANSIENT	CLUSTERING_KEY	ROW_COUNT	BYTES	RETENTION
1	TIPS_DATABASE	LAB	DEVICE_DATA	SYSADMIN	BASE TABLE	NO		0	0	
2	TIPS_DATABASE	LAB	DEVICE_DATA_CREATE_LIKE	SYSADMIN	BASE TABLE	NO		0	0	
3	TIPS_DATABASE	LAB	DEVICE_DATA_VARIANT	SYSADMIN	BASE TABLE	NO		1	30720	
4	TIPS_DATABASE	LAB	DEVICE_DATA_VARIANT_2	SYSADMIN	BASE TABLE	NO		1000	21504	
5	TIPS_DATABASE	LAB	ORC_CLONED_TBL	SYSADMIN	BASE TABLE	NO		19980	377856	
6	TIPS_DATABASE	LAB	ORC_DEVICE_DATA	SYSADMIN	BASE TABLE	NO		0	0	
7	TIPS_DATABASE	LAB	ORC_DEVICE_DATA_VARIANT	SYSADMIN	BASE TABLE	NO		19980	377856	
8	TIPS_DATABASE	LAB	PARQUET_DEVICE_DATA_AR...	ACCOUNTAD...	BASE TABLE	NO		0	0	
9	TIPS_DATABASE	LAB	PARQUET_DEVIDE_DATA_DB...	ACCOUNTAD...	BASE TABLE	NO		0	0	
10	TIPS_DATABASE	LAB	PARQUET_DEVIDE_DATA_VA...	SYSADMIN	BASE TABLE	NO		19980	377856	
11	TIPS_DATABASE	LAB	PARQUET_TYPE_ARRAY	SYSADMIN	BASE TABLE	NO		0	0	

Databases > TIPS_DATABASE_QA

Tables Views Schemas Stages File Formats Sequences Pipes

+ Create... + Create Like... Clone... Load Data... Drop... Transfer Ownership

Table Name	Schema	Creation Time	Owner	Rows	Size	Comment
ORC_CLONED_TBL	LAB	2:25:07 PM	SYSADMIN	20.0K	369KB	This table will have only o...
TESTING1	LAB	2:24:27 PM	SYSADMIN			test
DEVICE_DATA_CREATE_LIKE	LAB	2:23:32 PM	SYSADMIN			Create like feature is used ...
DEVICE_DATA	LAB	2:15:59 PM	SYSADMIN			This is customer iot devic...
DEVICE_DATA_VARIANT	LAB	2:15:59 PM	SYSADMIN	1	30KB	This is single column data
DEVICE_DATA_VARIANT_2	LAB	2:15:59 PM	SYSADMIN	1K	21KB	Trying to load on individua...
ORC_DEVICE_DATA	LAB	2:15:59 PM	SYSADMIN			Storing Data from ORC File
ORC_DEVICE_DATA_VARIANT	LAB	2:15:59 PM	SYSADMIN	20.0K	369KB	This table will have only o...
PARQUET DEVICE DATA ARRAY	LAB	2:15:59 PM	ACCOUNTADMIN			This is an experiment with...

Create Table Like

Name *

Source

Comment

Show SQL Cancel Finish

Databases > TIPS_DATABASE_QA

Tables Views Schemas Stages File Formats Sequences Pipes

+ Create... + Create Like... Clone... Load Data... Drop... Transfer Ownership

Table Name	Schema	Creation Time	Owner	Rows	Size	Comment
DEVICE_DATA	LAB	2:15:59 PM	SYSADMIN			This is customer iot device data
DEVICE_DATA_VARIANT	LAB	2:15:59 PM	SYSADMIN	1	30KB	This is single column data
DEVICE_DATA_VARIANT_2	LAB	2:15:59 PM	SYSADMIN	1K	21KB	Trying to load on individual rows
ORC_DEVICE_DATA	LAB	2:15:59 PM	SYSADMIN			Storing Data from ORC File
ORC_DEVICE_DATA_VARIANT	LAB	2:15:59 PM	SYSADMIN	20.0K	369KB	This table will have only one column
PARQUET_DEVICE_DATA_ARRAY	LAB	2:15:59 PM	ACCOUNTADMIN			This is an experiment with Parquet with array
PARQUET_DEVIDE_DATA_OBJ...	LAB	2:15:59 PM	ACCOUNTADMIN			This table has one field with object as field type
PARQUET_DEVIDE_DATA_VARI...	LAB	2:15:59 PM	SYSADMIN	20.0K	369KB	The single column data loaded via Variant
PARQUET_TYPE_ARRAY	LAB	2:15:59 PM	SYSADMIN			Parquet data loaded in array data type
TABLEAU_JSON	LAB	2:15:59 PM	SYSADMIN	5.5K	260KB	Sample JSON Data
DEMO_TABLE	PUBLIC	2:15:59 PM	SYSADMIN	2	1KB	
TEST_SEMI_STRUCTURED	PUBLIC	2:15:59 PM	SYSADMIN			
VARIA	PUBLIC	2:15:59 PM	SYSADMIN	6	1KB	

Database	Origin	↓ Creation Time	Owner	Comment
TIPS_DATABASE_QA		2:15 PM	SYSADMIN	This is practice Database

Clone Database

Name * TIPS_DATABASE_QA

Source TIPS_DATABASE

Comment This is cloned database from Tips Database

[Show SQL](#) Cancel Finish

Account Usage

- SNOWFLAKE is a system-defined, read-only shared database, provided by Snowflake
 - The database is automatically imported into each account from a share named ACCOUNT_USAGE
 - The SNOWFLAKE database is an example of Snowflake utilizing Secure Data Sharing to provide object metadata and other usage metrics for your account.
 - The SNOWFLAKE database contains three schemas (also read-only). Each schema contains a set of views (refer images)
-
- ACCOUNT_USAGE: Views that display object metadata and usage metrics for your account.
 - Diff in ACCOUNT_USAGE vs INFORMATION_SCHEMA
 - Records for dropped objects included in each view.
 - Longer retention time for historical usage data.
 - Data latency.
 - By default, only account administrators (users with the ACCOUNTADMIN role) can access the SNOWFLAKE database and schemas within the database, or perform queries on the views

-- Enabling Account Usage for Other Roles

```
USE ROLE ACCOUNTADMIN;
```

```
GRANT IMPORTED PRIVILEGES ON DATABASE snowflake TO ROLE SYSADMIN;
```

```
GRANT IMPORTED PRIVILEGES ON DATABASE snowflake TO ROLE customrole1;
```

```
USE ROLE customrole1;
```

```
SELECT * FROM snowflake.account_usage.databases;
```

```
-- Sample Queries
```

```
-- Average number of seconds between failed login attempts by user (month-to-date):
```

```
-- refer image for result and cost of the query
```

```
USE ROLE ACCOUNTADMIN;
```

```
select user_name,  
       count(*) as failed_logins,  
       avg(seconds_between_login_attempts) as average_seconds_between_login_attempts  
from (  
    select user_name,  
           timediff(seconds, event_timestamp, lead(event_timestamp)  
                  over(partition by user_name order by event_timestamp)) as  
           seconds_between_login_attempts  
    from "SNOWFLAKE"."ACCOUNT_USAGE"."LOGIN_HISTORY"  
   where event_timestamp > date_trunc(month, current_date)  
     and is_success = 'NO'  
 )  
group by 1  
order by 3;
```

```
--Failed logins by user (month-to-date):
```

```
-- refer image for result
```

```
select user_name,  
       sum(if(is_success = 'NO', 1, 0)) as failed_logins,  
       count(*) as logins,  
       sum(if(is_success = 'NO', 1, 0)) / nullif(count(*), 0) as login_failure_rate  
from "SNOWFLAKE"."ACCOUNT_USAGE"."LOGIN_HISTORY"  
where event_timestamp > date_trunc(month, current_date)  
group by 1  
order by 4 desc;
```

```
-- Failed logins by user and connecting client (month-to-date):
select reported_client_type,
       user_name,
       sum(if(is_success = 'NO', 1, 0)) as failed_logins,
       count(*) as logins,
       sum(if(is_success = 'NO', 1, 0)) / nullif(count(*), 0) as login_failure_rate
from "SNOWFLAKE"."ACCOUNT_USAGE"."LOGIN_HISTORY"
where event_timestamp > date_trunc(month, current_date)
group by 1,2
order by 5 desc;
```

-- Examples: Warehouse Credit Usage

-- Credits used by each warehouse in your account (month-to-date):

-- refer image for result

```
select warehouse_name,
       sum(credits_used) as total_credits_used
from "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_METERING_HISTORY"
where start_time >= date_trunc(month, current_date)
group by 1
order by 2 desc;
```

-- Credits used over time by each warehouse in your account (month-to-date):

-- refer image for result

```
select start_time::date as usage_date,
       warehouse_name,
       sum(credits_used) as total_credits_used
from "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_METERING_HISTORY"
where start_time >= date_trunc(month, current_date)
group by 1,2
```

```
order by 2,1;
```

-- Examples: Data Storage Usage

-- Billable terabytes stored in your account over time:

```
select date_trunc(month, usage_date) as usage_month  
     , avg(storage_bytes + stage_bytes + failsafe_bytes) / power(1024, 4) as billable_tb  
  from "SNOWFLAKE"."ACCOUNT_USAGE"."STORAGE_USAGE"  
 group by 1  
order by 1;
```

-- Examples: User Query Totals and Execution Times

-- Total jobs executed in your account (month-to-date):

```
select count(*) as number_of_jobs  
  from query_history  
 where start_time >= date_trunc(month, current_date);
```

-- Total jobs executed by each warehouse in your account (month-to-date):

```
select warehouse_name,  
       count(*) as number_of_jobs  
  from query_history  
 where start_time >= date_trunc(month, current_date)  
 group by 1  
order by 2 desc;
```

--Average query execution time by user (month-to-date):

```
select user_name,  
       avg(execution_time) as average_execution_time  
  from query_history  
 where start_time >= date_trunc(month, current_date)  
 group by 1  
 order by 2 desc;
```

--Average query execution time by query type and warehouse size (month-to-date):

```
select query_type,  
       warehouse_size,  
       avg(execution_time) as average_execution_time  
  from query_history  
 where start_time >= date_trunc(month, current_date)  
 group by 1,2  
 order by 3 desc;
```

Find database objects Starting with... Run All Queries Saved 4 minutes ago

METERING_DAILY_HISTORY
METERING_DAILY_HISTORY
PINES
PIPE_USAGE_HISTORY
QUERY_HISTORY
REFERENTIAL_CONSTRAINTS
REPLICATION_USAGE_HISTORY
ROLES
SCHEMATA
SEQUENCES
STAGES
STAGE_STORAGE_USAGE_HISTORY
STORAGE_USAGE
TABLES
TABLE_CONSTRAINTS
METERING_DAILY...

Results Data Preview View: SNOWFLAKE.ACCOUNT_USAGE.METERING_DAILY_HISTORY Open History

Row	SERVICE_TYPE	USAGE_DATE	CREDITS_USED_COMPUTE	CREDITS_USED_CLOUD_SERVICES	CREDITS_USED	CREDITS_ADJUSTMENT_CLOUD_SERVICE	CREDITS_BILLED
1	WAREHOUSE...	2020-05-08	0.00000000	0.000007500	0.000007500	0.000000000	0.0000075000
2	WAREHOUSE...	2020-05-08	0.67972222	0.003016667	0.682738889	-0.0030166670	0.679722220
3	WAREHOUSE...	2020-05-10	0.373888889	0.001405556	0.375294445	-0.0014055560	0.3738888890
4	WAREHOUSE...	2020-05-04	1.211388889	0.000486111	1.211875000	-0.0004861110	1.2118888890
5	WAREHOUSE...	2020-05-11	0.225555556	0.001213056	0.226788612	-0.0012130560	0.2255555560
6	WAREHOUSE...	2020-05-05	0.172500000	0.000960000	0.173460000	-0.0009600000	0.1725000000
7	WAREHOUSE...	2020-05-02	0.575833333	0.001318056	0.577151389	-0.0013180560	0.5758333330
8	WAREHOUSE...	2020-05-12	0.572777778	0.001247778	0.574205556	-0.0012477780	0.5727777780

Find database objects Starting with... Run All Queries Saved 2 minutes ago

STAGES
STAGE_STORAGE_USAGE_HISTORY
STORAGE_USAGE
TABLES
TABLE_CONSTRAINTS
TABLE_STORAGE_METRICS
USERS
VIEWS
WAREHOUSE_LOAD_HISTORY
WAREHOUSE_METERING_HISTORY
INFORMATION_SCHEMA
ORGANIZATION_USAGE
STORAGE_USAGE

Results Data Preview View: SNOWFLAKE.ACCOUNT_USAGE.STORAGE_USAGE Open History

Row	USAGE_DATE	STORAGE_BYTES	STAGE_BYTES	FAILSAFE_BYTES
1	2020-05-07	1532.000000	2621.000000	0.000000
2	2020-05-06	1532.000000	2621.000000	0.000000
3	2020-05-09	984253.000000	2712498.000000	0.000000
4	2020-05-02	1534.000000	2443.000000	0.000000
5	2020-05-03	1534.000000	2469.000000	0.000000
6	2020-04-29	0.000000	0.000000	0.000000
7	2020-04-30	0.000000	0.000000	0.000000
8	2020-05-04	1533.000000	2553.000000	0.000000
9	2020-05-05	1533.000000	2620.000000	0.000000
10	2020-05-10	1408267.000000	2712957.000000	195067.000000
11	2020-05-01	1535.000000	2447.000000	0.000000
12	2020-05-08	332446.000000	650666.000000	0.000000

Find database objects Starting with... Run All Queries Saved 1 minute ago

TABLE_CONSTRAINTS
TABLE_STORAGE_METRICS
USERS
VIEWS
WAREHOUSE_LOAD_HISTORY
WAREHOUSE_METERING_HISTORY
INFORMATION_SCHEMA
ORGANIZATION_USAGE
READER_ACCOUNT_USAGE
SNOWFLAKE_SAMPLE_DATA
TIPS_DATABASE
UDCA_NUTRIENT_STORE
TABLE_STORAGE...

Results Data Preview View: SNOWFLAKE.ACCOUNT_USAGE.TABLE_STORAGE_METRICS Open History

Row	ID	TABLE_NAME	TABLE_SCHEMA	TABLE_CATALOG	TABLE_CATALOG	CLONE_GROUP	IS_TRANSIENT	ACTIVE_BYTES	TIME_TRAVEL_B
1	2	FD_GROUP	4 PUBLIC	6 USDA_NUTR...	2 NO	0	0	0	0
2	4	FD_GROUP_INGEST	4 PUBLIC	6 USDA_NUTR...	4 NO	1536	0	0	0
3	2056	PARQUET_DEVICE_DA...	6 LAB	7 TIPS_DATAB...	2056 NO	0	0	0	0
4	2054	PARQUET_DEVIDE_DA...	6 LAB	7 TIPS_DATAB...	2054 NO	0	0	0	0
5	2058	PARQUET_TYPE_ARRAY	6 LAB	7 TIPS_DATAB...	2058 NO	0	0	0	0
6	1030	ORC_DEVICE_DATA	6 LAB	7 TIPS_DATAB...	1030 NO	0	0	0	0
7	2060	TABLEAU_JSON	6 LAB	7 TIPS_DATAB...	2060 NO	266240	0	0	0
8	2052	PARQUET_DEVIDE_DA...	6 LAB	7 TIPS_DATAB...	2052 NO	377856	0	0	0
9	1032	ORC_DEVICE_DATA_V...	6 LAB	7 TIPS_DATAB...	1032 NO	377856	0	0	0
10	2050	DEVICE_DATA_VARIANT	6 LAB	7 TIPS_DATAB...	2050 NO	30720	0	0	0
11	1028	DEVICE_DATA_VARIA...	6 LAB	7 TIPS_DATAB...	1028 NO	21504	0	0	0
12	1026	DEVICE DATA	6 LAB	7 TIPS_DATAB...	1026 NO	0	0	0	0

Find database objects Run All Queries Saved 10 seconds ago

Starting with... Open History

- USERS
- VIEWS
- WAREHOUSE_LOAD_HISTORY
- WAREHOUSE_METERING_HISTORY
- INFORMATION_SCHEMA
- ORGANIZATION_USAGE
- READER_ACCOUNT_USAGE
- SNOWFLAKE_SAMPLE_DATA
- TIPS_DATABASE
- USDA_NUTRIENT_STDREF
- UTIL_DB

WAREHOUSE_MET... Preview Data X

Columns Data Type

START_TIME	TIMESTAMP_LTZ(9)
END_TIME	TIMESTAMP_LTZ(9)
WAREHOUSE_ID	NUMBER(38,0)
WAREHOUSE_NAME	VARCHAR(1677216)
CREDITS_USED	NUMBER(38,9)
CREDITS_USED_COMPUTE	NUMBER(38,9)
CREDITS_USED_CLOUD_SER...	NUMBER(38,9)

Row START_TIME END_TIME WAREHOUSE_ID WAREHOUSE_NAME CREDITS_USED CREDITS_USED_COMPUT CREDITS_USED_CLOUD_S...

Row	START_TIME	END_TIME	WAREHOUSE_ID	WAREHOUSE_NAME	CREDITS_USED	CREDITS_USED_COMPUT	CREDITS_USED_CLOUD_S...
1	2020-05-04 02:00:0...	2020-05-04 03:00:0...	0	CLOUD_SERVICES_O...	0.000017778	0.000000000	0.000017778
2	2020-05-10 23:00:0...	2020-05-11 00:00:0...	1	COMPUTE_WH	0.226768611	0.225555556	0.001213056
3	2020-05-09 07:00:0...	2020-05-09 08:00:0...	1	COMPUTE_WH	0.444354444	0.442222222	0.002132222
4	2020-05-11 22:00:0...	2020-05-11 23:00:0...	0	CLOUD_SERVICES_O...	0.000045000	0.000000000	0.000045000
5	2020-05-04 02:00:0...	2020-05-04 03:00:0...	4	ADFAASDF	0.266690000	0.266666667	0.000023333
6	2020-05-11 22:00:0...	2020-05-11 23:00:0...	1	COMPUTE_WH	0.471902222	0.470833333	0.001068889
7	2020-05-09 08:00:0...	2020-05-09 09:00:0...	1	COMPUTE_WH	0.000178333	0.000000000	0.000178333
8	2020-05-04 02:00:0...	2020-05-04 03:00:0...	1	COMPUTE_WH	0.089261944	0.089166667	0.000095278
9	2020-05-09 04:00:0...	2020-05-09 05:00:0...	1	COMPUTE_WH	0.000151111	0.000000000	0.000151111
10	2020-05-02 07:00:0...	2020-05-02 08:00:0...	1	COMPUTE_WH	0.442004722	0.441388889	0.000615833
11	2020-05-04 02:00:0...	2020-05-04 03:00:0...	3	TESTING	0.484466944	0.484444444	0.000022500
12	2020-05-02 07:00:0...	2020-05-02 08:00:0...	0	CLOUD_SERVICES_O...	0.000173056	0.000000000	0.000173056

Run All Queries Saved 13 seconds ago Open History

```
1 select date_trunc(month, usage_date) as usage_month
2 , avg(storage_bytes + stage_bytes + failsafe_bytes) / power(1024, 4) as billable_tb
3 from "SNOWFLAKE"."ACCOUNT_USAGE"."STORAGE_USAGE"
4 group by 1
5 order by 1;
```

Results Data Preview Open History

✓ Query ID SQL 895ms 2 rows

Filter result... Download Copy Columns

Row	USAGE_MONTH	BILLABLE_TB
1	2020-04-01	0
2	2020-05-01	1.10333291e-06

Run All Queries Saved 25 seconds ago Open History

```
1 select start_time::date as usage_date,
2      warehouse_name,
3      sum(credits_used) as total_credits_used
4  from "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_METERING_HISTORY"
5 where start_time >= date_trunc(month, current_date)
6 group by 1,2
7 order by 2,1;
```

Results Data Preview Open History

✓ Query ID SQL 1.11s 15 rows

Filter result... Download Copy Columns

Row	USAGE_DATE	WAREHOUSE_NAME	TOTAL_CREDITS_USED
1	2020-05-04	ADFAASDF	0.266690000
2	2020-05-02	CLOUD_SERVICES_ONLY	0.000702223
3	2020-05-04	CLOUD_SERVICES_ONLY	0.000017778
4	2020-05-05	CLOUD_SERVICES_ONLY	0.000006111
5	2020-05-11	CLOUD_SERVICES_ONLY	0.000045000
6	2020-05-04	COMPUTE_2CREDIT	0.371451111
7	2020-05-05	COMPUTE_2CREDIT	0.000005556
8	2020-05-02	COMPUTE_WH	0.576449166
9	2020-05-04	COMPUTE_WH	0.089261944

1 select warehouse_name,
2 sum(credits_used) as total_credits_used
3 from "SNOWFLAKE"."ACCOUNT_USAGE"."WAREHOUSE_METERING_HISTORY"
4 where start_time >= date_trunc(month, current_date)
5 group by 1
6 order by 2 desc;

Results Data Preview ← Open History

✓ Query ID SQL 1.36s 5 rows

Filter result...

Row	WAREHOUSE_NAME	TOTAL_CREDITS_USED
1	COMPUTE_WH	2.595891664
2	TESTING	0.484466944
3	COMPUTE_2CREDIT	0.371456667
4	ADFAASDF	0.266690000
5	CLOUD_SERVICES_ONLY	0.000771112

History

Add a filter to customize your results.

Include client-generated statements
 Include queries executed by user tasks

Status	Incident	Query ID	SQL Text	User	Warehouse	Size
✓		01942c60-...	select user_name, count(*) as failed_logins, avg(seconds_between_login_attempts) as average...	USER01	COMPUTE_...	X-Small
✓		01942c43-...	select warehouse_name, sum(credits_used) as total_credits_us...			
✓		01942c43-...	use role accountadmin;			
✗		01942c42-...	select warehouse_name, sum(credits_used) as total_credits_us...			
✓		01942c42-...	ALTER WAREHOUSE "COMPUTE_WH" RESUME;	USER01	COMPUTE_...	

Find database objects All Queries | Saved 1 minute ago

Starting with...

- DEMO_DB
- MAY022020DB
- SNOWFLAKE**
 - ACCOUNT_USAGE**
 - No Tables in this Schema
 - Views
 - AUTOMATIC_CLUSTERING_HISTORY
 - COLUMNS
 - COPY_HISTORY
 - DATABASES
 - DATABASE_STORAGE_USAGE_HIST...
 - DATA_TRANSFER_HISTORY
 - FILE_FORMATS
 - FUNCTIONS
 - GRANTS_TO_ROLES
 - GRANTS_TO_USERS
 - LOAD_HISTORY
 - LOGIN_HISTORY
 - MATERIALIZED_VIEW_REFRESH_HI...
 - METERING_DAILY_HISTORY
 - METERING_HISTORY

```

1 select user_name,
2      count(*) as failed_logins,
3      avg(seconds_between_login_attempts) as average_seconds_between_login_attempts
4  from (
5      select user_name,
6             timediff(seconds, event_timestamp, lead(event_timestamp))
7                 over(partition by user_name order by event_timestamp) as seconds_between_login_attempts
8      from "SNOWFLAKE"."ACCOUNT_USAGE"."LOGIN_HISTORY"
9      where event_timestamp > date_trunc(month, current_date)
10         and is_success = 'NO'
11    )
12   group by 1
13  order by 3;

```

Total Duration 1.43s

- Queuing 100ms
- Compilation 521ms
- Execution 805ms
- Compiling SQL 521ms
- Executing (serial) 2ms
- Queued (resuming warehouse partition) 100ms
- Wait process group 203ms
- Executing (warehouse) 534ms
- Receiving request from client 24ms

Results

✓ Query ID SQL 1.43s 0 rows

Filter result...

 ORGANIZATION_USAGE
No Tables in this Schema
▼ Views
🔗 PREVIEW_DATA_TRANSFER_DAILY_HIS...
🔗 PREVIEW_METERING_DAILY_HISTORY
🔗 PREVIEW_STORAGE_DAILY_HISTORY

 READER_ACCOUNT_USAGE
No Tables in this Schema
▼ Views
🔗 LOGIN_HISTORY
🔗 QUERY_HISTORY
🔗 RESOURCE_MONITORS
🔗 STORAGE_USAGE
🔗 WAREHOUSE_METERING_HISTORY

 SNOWFLAKE
 ACCOUNT_USAGE
No Tables in this Schema
▼ Views
🔗 AUTOMATIC_CLUSTERING_HISTORY
🔗 COLUMNS
🔗 COPY_HISTORY
🔗 DATABASES
🔗 DATABASE_STORAGE_USAGE_HISTO...
🔗 DATA_TRANSFER_HISTORY
🔗 FILE_FORMATS
🔗 FUNCTIONS
🔗 GRANTS_TO_ROLES
🔗 GRANTS_TO_USERS
🔗 LOAD_HISTORY

 SNOWFLAKE
 ACCOUNT_USAGE
 INFORMATION_SCHEMA
 ORGANIZATION_USAGE
 READER_ACCOUNT_USAGE

WebUI

The screenshot shows the Snowflake WebUI interface. At the top, there is a navigation bar with icons for Databases, Shares (selected), Data Marketplace, Warehouses, Worksheets, History, Account, Partner Connect, Help, and Notifications. On the far right, it shows the user 'USER01 ACCOUNTADMIN'. Below the navigation bar, the title 'Secure Shares' is displayed, along with a timestamp 'Last refreshed 9:37:14 AM' and a refresh button. There are tabs for 'Inbound', 'Outbound', and '+ Create', followed by a link to 'Create Database From Secure Share'. A search bar labeled 'Search Inbound Secure Shares' shows '2 Inbound Secure Shares'. A table lists the shares with columns: Secure Share Name, Shared By, Database, Creation Time, Owner, and Comment. The data shows two entries: 'ACCOUNT_USAGE' shared by 'SNOWFLAKE' in database 'SNOWFLAKE' at '9/19/2019, 5:12:31 AM', and 'SAMPLE_DATA' shared by 'SFC_SAMPLES' in database 'SNOWFLAKE_SAMPLE_DA...' at '8/20/2019, 5:59:20 AM'. A 'Columns' dropdown menu is visible on the right.

The screenshot shows the Snowflake WebUI with a modal window overlaid. The modal has a blue header with the text 'Discover new external data'. Below the header, there is a message: 'Drive business insights, instantly, with data from the world's best data and software providers on the Snowflake Data Marketplace.' A button labeled 'Explore the Snowflake Data Marketplace →' is present. Below the button, it says 'Now, on the new Snowflake Experience'. The background of the main interface shows the same navigation bar and 'Shares' page content as the first screenshot, but they are dimmed due to the modal being active.

The screenshot shows the Snowflake Worksheet interface. At the top, there are navigation tabs: Databases, Shares, Data Marketplace, Warehouses, Worksheets (which is selected), History, and Account. On the right, there's a user profile for 'USER01 SECURITYADMIN'. Below the tabs, a search bar says 'variant' and 'Find database objects Starting with...'. A code editor window contains the following JSON query:

```

1 select
2   vm.value::string as make,
3   vm.value::string as model,
4   ve.value::string as "Extras Purchased"
5   from
6   car_sales
7   , lateral flatten(input => src:vehicle) vm
8   , lateral flatten(input => vm.value:extras) ve;

```

The results section shows a single row of data:

Row	created_on	role	granted_to	grantee_name	granted_by
1	2020-05-02 ...	ACCOUNTA...	USER	USER01	

Below the results is a history table with 21 entries. The columns are: Status, Duration, Start, End, Query ID, Rows, and Bytes Scanned.

This screenshot shows the Snowflake Worksheet interface with several callout boxes highlighting various features:

- Add a worksheet**: Points to the '+' button in the top right of the worksheet area.
- SQL editor**: Points to the SQL code editor area.
- Object browser**: Points to the sidebar where databases like ABC_EMPLOYEES, HOLIDAYS, HR, EMP_ADDR, and EMP_PH are listed.
- Manage worksheets (search for, open or delete, rename)**: Points to the top right of the worksheet area.
- Open a tutorial**: Points to the top right of the worksheet area.
- Load a script**: Points to the top right of the worksheet area.
- Change the current database, schema, or warehouse for the current worksheet without losing your work.**: Points to the context menu at the top right.
- Resume/suspend or resize your current warehouse.**: Points to the context menu at the top right.
- Hide or show columns**: Points to the 'Columns' dropdown in the results table.
- Download results**: Points to the download icon in the results table toolbar.
- Copy results to clipboard**: Points to the copy icon in the results table toolbar.
- Maximize/restore results**: Points to the maximize/minimize button in the results table header.

The worksheet area displays a query to select data from the emp table, showing rows 51 to 58. The results table shows three rows of employee data with columns: Row, ID, FIRST_NAME, LAST_NAME, CITY, and POSTAL_CODE.

User Role Creation

--create a user (Blr@4950\$)

```
CREATE USER tips_jr_analyst
```

```
PASSWORD = '*****'
```

```
COMMENT = 'Tips Junior Analyst to run queries'
```

```
MUST_CHANGE_PASSWORD = TRUE;
```

-- with advance parameters

```
CREATE USER tips_jr_analyst
```

```
PASSWORD = '*****'
```

```
COMMENT = 'Tips Junior Analyst to run queries'
```

```
LOGIN_NAME = 'TipsJrAnalyst'
```

```
DISPLAY_NAME = 'Junior Analyst'
```

```
FIRST_NAME = 'Junior'
```

```
LAST_NAME = 'Analyst'
```

```
EMAIL = 'tips.analyst@toppertips.com'
```

```
MUST_CHANGE_PASSWORD = TRUE;
```

-- With advance params & default warehouse details

-- also granting role

```
CREATE USER tips_jr_analyst
```

```
PASSWORD = '*****'
```

```
COMMENT = 'Tips Junior Analyst to run queries'
```

```
LOGIN_NAME = 'TipsJrAnalyst'
```

```
DISPLAY_NAME = 'Junior Analyst'
```

```
FIRST_NAME = 'Junior'
```

```
LAST_NAME = 'Analyst'
```

```
EMAIL = 'tips.analyst@toppertips.com'
```

```
DEFAULT_ROLE = "USERADMIN"  
DEFAULT_WAREHOUSE = 'COMPUTE_WH'  
DEFAULT_NAMESPACE = 'TIPS_DB'  
MUST_CHANGE_PASSWORD = TRUE;
```

```
GRANT ROLE "USERADMIN" TO USER tips_jr_analyst;
```

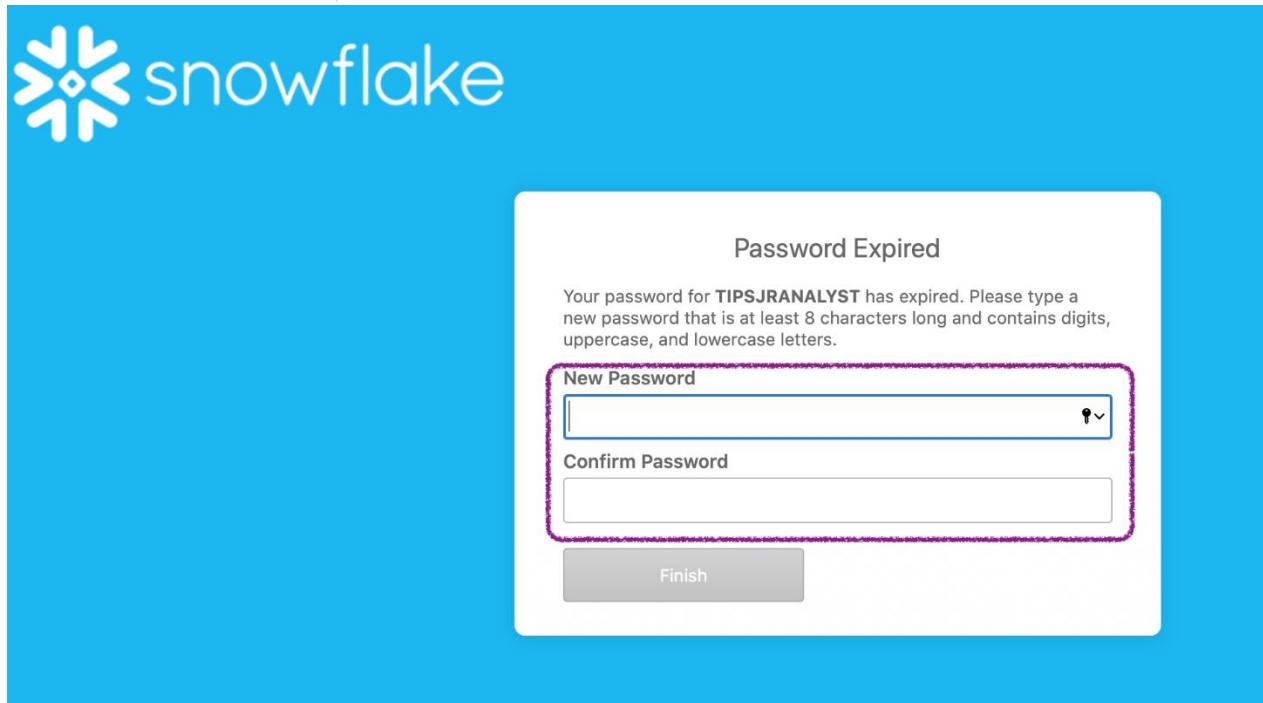
```
-- for user to see the warehouse, it has to be assigned to that role
```

```
GRANT USAGE, OPERATE, MONITOR ON WAREHOUSE "COMPUTE_WH"  
TO ROLE "USERADMIN";
```

```
-- grant the database usage
```

```
GRANT USAGE, CREATE SCHEMA, MODIFY, MONITOR ON DATABASE "TIPS_SALES_PROD"  
TO ROLE "USERADMIN";
```

The screenshot shows the Snowflake Worksheets interface. At the top, there are navigation icons for Databases, Shares, Warehouses, Worksheets (which is selected), and History. On the right, there are links for Partner Connect, Help, and a user account labeled "Junior Analyst USERADMIN". A context menu is open over a database object named "TIPS_SALES_PROD" in the left sidebar, which is highlighted with a red box. The context menu includes options for Role (USERADMIN), Warehouse (COMPUTE_WH XS Suspended), Database (TIPS_SALES_PROD), and Schema (INFORMATION_SCHEMA). The main workspace is empty, showing a placeholder message: "Query results will appear here."



Snowflake Account Overview

Navigation: Databases, Shares, Warehouses, Worksheets, History, Account (selected)

User Management:

- Usage, Billing, **Users**, Roles, Policies, Sessions, Resource Monitors, Reader Accounts

Create User: [+ Create...](#)

User Name	Status	Last Login	Time Left	Login Name	Display Name	Creation Time
SNOWFLAKE	Expired	5/2/2020, 7:04:21 PM		SNOWFLAKE	SNOWFLAKE	5/2/2020, 7:04:21 PM
TIPS_JR_ANALYST	Enabled			TIPSJRANALYST	Junior Analyst	1:37:44 PM
USER01	Enabled	1:29:42 PM		USER01	USER01	5/2/2020, 7:04:21 PM

Detailed View for TIPS_JR_ANALYST:

- Disable User Reset Password Drop
- Login Name: TIPSJRANALYST
- Display Name: Junior Analyst
- Default Role: USERADMIN
- Default Warehouse: COMPUTE_WH
- Default Namespace: TIPS_DB
- MFA
- Last Login
- Status: Enabled

Created on 5/16/2020 at 1:37:44 PM

Create User

General Advanced Preferences

User Name * tips_jr_analyst

New Password *

Confirm Password *

Comment: Tips Junior Analyst to run queries

Force Password Change [?](#)

Show SQL Cancel Next **Finish**

Grant SQL Example

-- Giving a usage grant on a table to public role

```
GRANT USAGE ON DATABASE "DB_FOR_SECURITY" TO ROLE "PUBLIC" WITH GRANT OPTION;
```

-- How To: Grant a role access to database objects in a schema

To allow a role to use database objects in a specific schema, the owner of the database objects (typically a system administrator (SYSADMIN role)) must grant privileges on the database, schema, and objects.

--Grant usage on the database:

```
GRANT USAGE ON DATABASE <database> TO ROLE <role>;
```

--Grant usage on the schema:

```
GRANT USAGE ON SCHEMA <database>.<schema> TO ROLE <role>;
```

--Grant the ability to query an existing table:

```
GRANT SELECT ON TABLE <database>.<schema>.<table> TO ROLE <role>;
```

The following table privileges are supported:

Privilege	Usage
SELECT	Execute a SELECT statement on the table.
INSERT	Execute an INSERT command on the table.
UPDATE	Execute an UPDATE command on the table.
TRUNCATE	Execute a TRUNCATE command on the table.
DELETE	Execute a DELETE command on the table.
REFERENCES	Reference the table as the unique/primary key table for a foreign key constraint.
ALL [PRIVILEGES]	Grant all privileges, except OWNERSHIP, on the table.
OWNERSHIP	Grant full control over a table.

Snowflake Name Stages

-- How to show all the stages?

```
show stages;  
  
list @S3_USDF_EXTERNAL_NAMED_STAGE;          --loading all files inside the bucket  
  
list @S3_USDF_EXTERNAL_NAMED_STAGE/load;      -- loading specific directory  
  
list @S3_USDF_EXTERNAL_NAMED_STAGE/load/D;    -- loading specific directory and file within it  
(look for * is missing)
```

-- Note : it is not stage, it is stages, make sure you remember the command.

-- Create a name internal stage (a temporary loading area)

-- This stage will be created within public schema if there is no other schema.

-- you can change the schema and databases names

```
CREATE OR REPLACE STAGE "USDA_NUTRIENT_STDREF"."PUBLIC".Tips_Internal_Storage  
COMMENT = 'An internal staging area';
```

-- The grant query to be documented?

-- Create a named external stage (a temporary loading area) with AWS S3 (Simple Storage Services)

-- This stage will be created within public schema if there is no other schema.

-- you can change the schema and databases names

```
CREATE OR REPLACE STAGE "USDA_NUTRIENT_STDREF"."PUBLIC".S3_Tips_External_Named_Stage  
URL = 's3://my-sample-bucket'  
CREDENTIALS = (AWS_KEY_ID = 'myuser' AWS_SECRET_KEY = '*****')  
ENCRYPTION = (MASTER_KEY = '*****')  
COMMENT = 'S3_Tips_External_Named_Stage';
```

-- Creating a temporary storage

-- Not documentation is found how it is different from normal stage and what is the purpose

-- This is to be explored ???

CREATE OR REPLACE TEMPORARY STAGE

"USDA_NUTRIENT_STDREF"."PUBLIC".S3_TEMPO_EXTERNAL_NAMED_STAGE

URL = 's3://my-sample-bucket'

CREDENTIALS = (AWS_KEY_ID = 'myuser' AWS_SECRET_KEY = '*****')

COMMENT = 'Checking what is this temporary stage is';

-- Alter Stage (Changing Parameters)

ALTER STAGE "USDA_NUTRIENT_STDREF"."PUBLIC".S3_USDF_EXTERNAL_NAMED_STAGE"

SET URL = 's3://on-demand-files'

COMMENT = 'S3_USDF_EXTERNAL_NAMED_STAGS';

-- Alter rename does not work, looks like it is not supported as UI also does not allow.

ALTER STAGE IF EXISTS

"USDA_NUTRIENT_STDREF"."PUBLIC".S3_TEMPO_EXTERNAL_NAMED_STAGE

RENAME "USDA_NUTRIENT_STDREF"."PUBLIC".S3_TEMPO_EXTERNAL_NAMED_STAGE_RENAMED

-- (this will end with error : SQL compilation error: syntax error line 2 at position 7 unexpected

"USDA_NUTRIENT_STDREF".)

-- Even unset comment does not seem to be working

-- (error : Unsupported feature 'UNSET').)

-- Note-01: if you input incorrect AWS keys, Snowflake will throw error like this

--Unable to create stage "S3_TIPS_EXTERNAL_NAMED_STAGE".

--The provided master key has invalid length. It must be either 128 bits, 192 bits, or 256 bits long.

-- Note-02: If the bucket(s3) is publicly available, then you don't need to specify the key at all.

-- Creating External Named Stages with Azure

```
CREATE OR REPLACE STAGE
"USDA_NUTRIENT_STDREF"."PUBLIC".Azure_Tips_External_Named_Stage
URL = 'azure://azurebucket'
CREDENTIALS = (AZURE_SAS_TOKEN = '*****')
ENCRYPTION = (TYPE = 'AZURE_CSE' MASTER_KEY =
'*****')
COMMENT = 'Azure_Tips_External_Named_Stage';
```

-- How to create a stage by cloning it

-- Creating a clone works for external stage but does not work for internal stages

```
CREATE STAGE "USDA_NUTRIENT_STDREF"."PUBLIC".Cloned_s3_tips
CLONE "USDA_NUTRIENT_STDREF"."PUBLIC"."S3_TIPS_EXTERNAL_NAMED_STAGE"
COMMENT = 'Cloning an s3 named stage';
```

-- Note-01 : If you try to create a clone using internal stages, it will throw error

Unable to create stage "CLONE_INTERNAL_STAGE".

Unsupported feature 'Cloning internal and temporary stages'.

-- Note-02 : Grants are also not copied when you clone

Different kind of error may appear if you are not having sufficient privileges

so make sure you check the role using the role manu and then run the queries or operations.

Check for the ICON called "Owner" on WebUI before operating on any staged object

Unable to modify stage "S3_TIPS_EXTERNAL_NAMED_STAGE".

SQL access control error: Insufficient privileges to operate on stage
'S3_TIPS_EXTERNAL_NAMED_STAGE'

Copy

/option_	owner	comment	region	type	cloud	notification_cha
	SYSADMIN	CLONED_S3...	us-east-1	EXTERNAL	AWS	NULL
	SYSADMIN	Checking w...	us-west-2	EXTERNAL TEMPORARY	AWS	NULL
	ACCOUNTA...	S3_Tips_Ext...	us-east-1	EXTERNAL	AWS	NULL
	SYSADMIN	S3_USDF_EX...	us-west-2	EXTERNAL	AWS	NULL
	ACCOUNTA...	An internal s...	NULL	INTERNAL	NULL	NULL

Results Data Preview

View: USDA_NUTRIENT_STDREF.INFORMATION_SCHEMA.STAGES

Filter result...

Row	STAGE_CATALOG	STAGE_SCHEMA	STAGE_NAME	STAGE_URL	STAGE_REGION	STAGE_TYPE	STAGE_OWNER
1	USDA_NUTRIENT_STDREF	PUBLIC	CLONED_S3_USDF_EXTER...	s3://my-sam...	us-east-1	External Named	SYSADMIN
2	USDA_NUTRIENT_STDREF	PUBLIC	S3_TIPS_EXTERNAL_NAM...	s3://my-sam...	us-east-1	External Named	ACCOUNTADMIN
3	USDA_NUTRIENT_STDREF	PUBLIC	S3_USDF_EXTERNAL_NA...	s3://on-demand...	us-west-2	External Named	SYSADMIN
4	USDA_NUTRIENT_STDREF	PUBLIC	TIPS_INTERNAL_STORAGE			Internal Named	ACCOUNTADMIN

STAGES

Preview Data X

Columns	Data Type
STAGE_CATALOG	VARCHAR(16777216)
STAGE_SCHEMA	VARCHAR(16777216)
STAGE_NAME	VARCHAR(16777216)
STAGE_URL	VARCHAR(16777216)
STAGE_REGION	VARCHAR(16777216)
STAGE_TYPE	VARCHAR(16777216)
STAGE_OWNER	VARCHAR(16777216)
COMMENT	VARCHAR(16777216)
CREATED	TIMESTAMP_LTZ(9)
LAST_ALTERED	TIMESTAMP_LTZ(9)

Databases > USDA_NUTRIENT_STDREF

Last refreshed 2:01:28 PM

Tables Views Schemas Stages File Formats Sequences Pipes

+ Create... Clone... Edit... Drop... Transfer Ownership

Stage	Schema	Location	Creation Time
CLONED_S3_USDF_EXTERNAL...	PUBLIC	s3://my-sample-bucket	2:01:23 PM
S3_USDF_EXTERNAL_NAMED_...	PUBLIC	s3://on-demand-files	2:01:06 PM
S3_TIPS_EXTERNAL_NAMED_...	Stage: S3_USDF_EXTERNAL_NAMED_STAGE		1:29:14 PM
TIPS_INTERNAL_STORAGE	PUBLIC	Snowflake	1:20:22 PM

S3_USDF_EXTERNAL_NAMED_STAGE

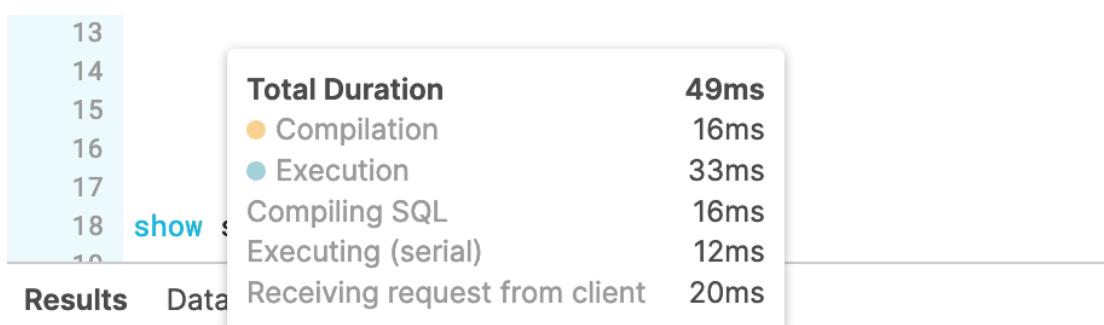
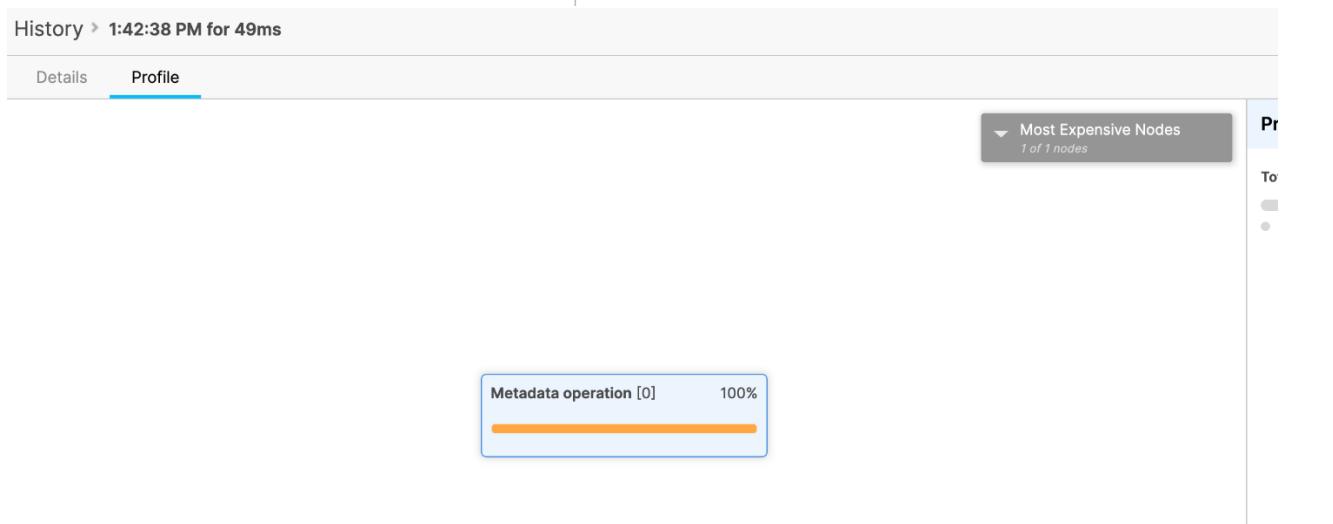
Owner

+ Grant Privileges

No Permissions Granted

History > 1:42:38 PM for 49ms

Details	Profile					
Status	Success					
User	USER01					
Warehouse	COMPUTE_WH					
Start Time	1:42:38 PM					
End Time	1:42:38 PM					
Total Duration	49ms					
Scanned Bytes	0					
Rows	0					
Query ID	019404cc-0370-410d-0000-001a208381fd					
Session ID	112214622217					
SQL Text						
1 show stages;						
Query Result						
Results						
row#	created_on	name	database_n...	schema_na...	url	has_
1	2020-05-05...	S3_TIPS_EX...	USDA_NUTR...	PUBLIC	s3://my-sam...	Y
2	2020-05-05...	TIPS_INTER...	USDA_NUTR...	PUBLIC		N



✓ Query_ID SQL 49ms 2 rows

Filter result...



Copy

Row	created_on	name	database_name	schema_nar
-----	------------	------	---------------	------------

18 show stages;

Results Data Preview

✓ Query ID SQL 49ms 2 rows

Filter result...

Row	created_on	name	database_name	schema_name	url	has_credentials	has_encryption	owner	comment	region
1	2020-05-05...	S3_TIPS_EX...	USDA_NUTR...	PUBLIC	s3://my-sam...	Y	N	ACCOUNTA...	S3_Tips_Ext...	us-east-1
2	2020-05-05...	TIPS_INTER...	USDA_NUTR...	PUBLIC		N	N	ACCOUNTA...	An internal s...	NULL

Clone Stage

Name *

Source

Comment

Create Stage

Staged files will be stored in the specified Azure location

Name *

Schema Name

URL *

Azure SAS Token

Encryption Master Key

Comment

Create Stage

Staged files will be stored in the specified S3 location

Name *

Schema Name

URL *

AWS Key ID

AWS Secret Key

Encryption Master Key

Comment

Edit privileges on stage TIPS_INTERNAL_STORAGE

Privileges to grant: READ, WRITE

Grant privileges to: PUBLIC

with Grant Option

Revoke All Cancel Grant

Grant privileges on stage TIPS_INTERNAL_STORAGE

Privileges to grant: READ, WRITE

Grant privileges to: Select a role

with Grant Option

ACCOUNTADMIN
PUBLIC
SECURITYADMIN
SYSADMIN
USERADMIN

Grant

Create Stage
Staged files will be stored in a Snowflake managed stage

Name*:

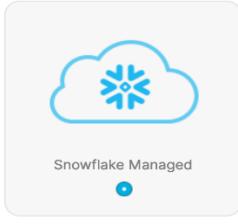
Schema Name: PUBLIC

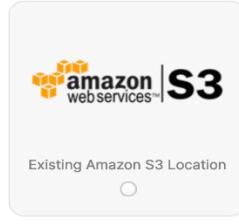
Comment: PUBLIC

Show SQL Cancel Back Finish

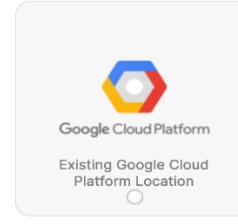
Create Stage

Choose a location for files to be staged

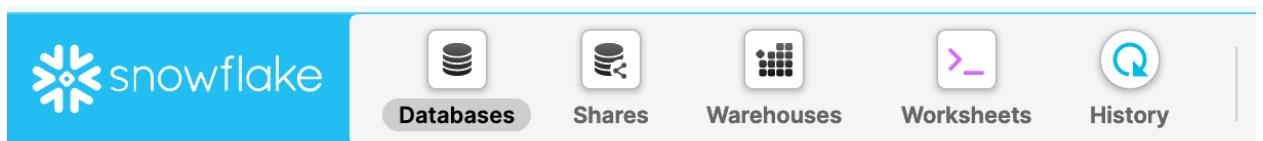
 Snowflake Managed

 Existing Amazon S3 Location

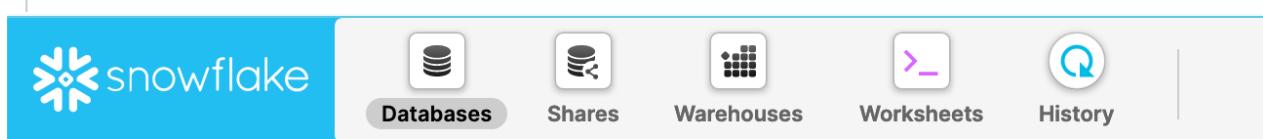
 Existing Microsoft Azure Location

 Existing Google Cloud Platform Location

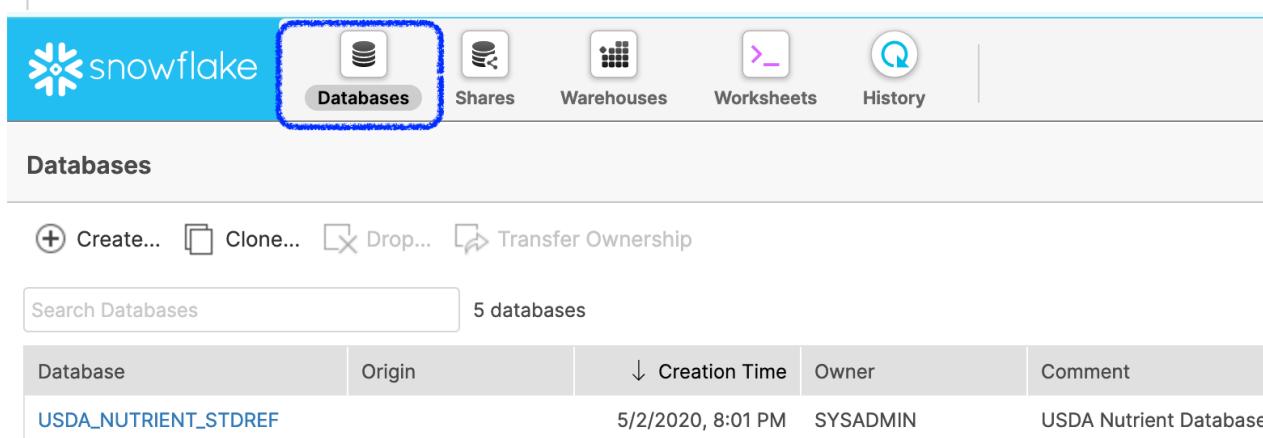
Cancel Next



The screenshot shows the Snowflake web interface. At the top, there's a blue header bar with the Snowflake logo on the left and five navigation icons on the right: Databases (selected), Shares, Warehouses, Worksheets, and History. Below the header, the title "Databases > USDA_NUTRIENT_STDREF" is displayed. A horizontal menu bar follows, with tabs for Tables, Views, Schemas, Stages (selected), File Formats, Sequences, and Pipes. Underneath the menu, there are several action buttons: "+ Create...", "Clone...", "Edit...", "Drop...", and "Transfer Ownership".



This screenshot shows the "Stages" page within the same database context. The top navigation bar and menu bar are identical to the previous screenshot. The title "Databases > USDA_NUTRIENT_STDREF" is also present. The "Tables" tab is selected in the menu bar. Below the menu, there are action buttons: "+ Create...", "+ Create Like...", "Clone...", "Load Data...", "Drop...", and "Transfer Ownersh".



This screenshot shows the "Databases" page. The top navigation bar and menu bar are identical to the previous screenshots. The title "Databases" is displayed. The "Tables" tab is selected in the menu bar. Below the menu, there are action buttons: "+ Create...", "Clone...", "Drop...", and "Transfer Ownership". A search bar labeled "Search Databases" is present, showing "5 databases". A table below lists the databases, with one row highlighted in blue.

Database	Origin	Creation Time	Owner	Comment
USDA_NUTRIENT_STDREF		5/2/2020, 8:01 PM	SYSADMIN	USDA Nutrient Database

Role

Important Notes

- For any role to function, we need at least one user assigned to it
- SYSADMIN role by default cannot create new roles or users. role creation operation will fail

```
create role tips_jr_dba;  
grant role tips_jr_dba to user tips_dba_user02;  
use role tips_jr_dba;
```

-- Now we have to give usage access to this new role

```
use role accountadmin;  
grant usage on database citibike to role tips_jr_dba;  
grant usage on database weather to role tips_jr_dba;
```

Hive

Hive Function Meta commands

- SHOW FUNCTIONS– lists Hive functions and operators
- DESCRIBE FUNCTION [function name]– displays short description of the function
- DESCRIBE FUNCTION EXTENDED [function name]– access extended description of the function

Types of Hive Functions

- UDF– is a function that takes one or more columns from a row as argument and returns a single value or object. Eg: concat(col1, col2)
- UDAF- aggregates column values in multiple rows and returns a single value. Eg: sum(c1)
- UDTF— takes zero or more inputs and produces multiple columns or rows of output. Eg: explode()
- Macros— a function that uses other Hive functions.

Mathematical Functions

Return Type	Name (Signature)	Description
BIGINT	round(double a)	Returns the rounded BIGINT value of the double
DOUBLE	round(double a, int d)	Returns the double rounded to d decimal places
BIGINT	floor(double a)	Returns the maximum BIGINT value that is equal or less than the double
BIGINT	ceil(double a), ceiling(double a)	Returns the minimum BIGINT value that is equal or greater than the double
double	rand(), rand(int seed)	Returns a random number (that changes from row to row) that is distributed uniformly from 0 to 1. Specifying the seed will make sure the generated random number sequence is deterministic.
double	exp(double a)	Returns ea where e is the base of the natural logarithm
double	ln(double a)	Returns the natural logarithm of the argument
double	log10(double a)	Returns the base-10 logarithm of the argument
double	log2(double a)	Returns the base-2 logarithm of the argument
double	log(double base, double a)	Return the base "base" logarithm of the argument
double	pow(double a, double p), power(double a, double p)	Return ap
double	sqrt(double a)	Returns the square root of a
string	bin(BIGINT a)	Returns the number in binary format
string	hex(BIGINT a) hex(string a)	If the argument is an int, hex returns the number as a string in hex format. Otherwise if the number is a string, it converts each character into its hex representation and returns the resulting string.
string	unhex(string a)	Inverse of hex. Interprets each pair of characters as a hexadecimal number and converts to the character represented by the number.
string	conv(BIGINT num, int from_base, int to_base), conv(STRING num, int from_base, int to_base)	Converts a number from a given base to another
double	abs(double a)	Returns the absolute value
int double	pmod(int a, int b) pmod(double a, double b)	Returns the positive value of a mod b
double	sin(double a)	Returns the sine of a (a is in radians)
double	asin(double a)	Returns the arc sin of x if -1<=a<=1 or null otherwise
double	cos(double a)	Returns the cosine of a (a is in radians)
double	acos(double a)	Returns the arc cosine of x if -1<=a<=1 or null otherwise
double	tan(double a)	Returns the tangent of a (a is in radians)
double	atan(double a)	Returns the arctangent of a
double	degrees(double a)	Converts value of a from radians to degrees
double	radians(double a)	Converts value of a from degrees to radians
int double	positive(int a), positive(double a)	Returns a
int double	negative(int a), negative(double a)	Returns -a
float	sign(double a)	Returns the sign of a as '1.0' or '-1.0'
double	e()	Returns the value of e
double	pi()	Returns the value of pi

Date Functions		
Return Type	Name (Signature)	Description
string	from_unixtime(bigint unixtime[, string format])	Converts the number of seconds from unix epoch (1970-01-01 00:00:00 UTC) to a string representing the timestamp of that moment in the current system time zone in the format of "1970-01-01 00:00:00"
bigint	unix_timestamp()	Gets current time stamp using the default time zone.
bigint	unix_timestamp(string date)	Converts time string in format yyyy-MM-dd HH:mm:ss to Unix time stamp, return 0 if fail; unix_timestamp('2009-03-20 11:30:01') = 1237573801
bigint	unix_timestamp(string date, string pattern)	Convert time string with given pattern to Unix time stamp, return 0 if fail: unix_timestamp('2009-03-20', 'yyyy-MM-dd') = 1237532400
string int	to_date(string timestamp) year(string date)	Returns the date part of a timestamp string: to_date("1970-01-01 00:00:00") = "1970-01-01" Returns the year part of a date or a timestamp string: year("1970-01-01 00:00:00") = 1970, year("1970-01-01") = 1970
int	month(string date)	Returns the month part of a date or a timestamp string: month("1970-11-01 00:00:00") = 11, month("1970-11-01") = 11
int int int int int	day(string date) dayofmonth(date) hour(string date) minute(string date) second(string date) weekofyear(string date)	Return the day part of a date or a timestamp string: day("1970-11-01 00:00:00") = 1, day("1970-11-01") = 1 Returns the hour of the timestamp: hour('2009-07-30 12:58:59') = 12, hour('12:58:59') = 12 Returns the minute of the timestamp Returns the second of the timestamp Return the week number of a timestamp string: weekofyear("1970-11-01 00:00:00") = 44, weekofyear("1970-11-01") = 44
int string string timestamp timestamp	datediff(string enddate, string startdate) date_add(string startdate, int days) date_sub(string startdate, int days) from_utc_timestamp(timestamp, string timezone) to_utc_timestamp(timestamp, string timezone)	Return the number of days from startdate to enddate: datediff('2009-03-01','2009-02-27') = 2 Add a number of days to startdate: date_add('2008-12-31', 1) = '2009-01-01' Subtract a number of days to startdate: date_sub('2008-12-31', 1) = '2008-12-30' Assumes given timestamp is UTC and converts to given timezone (as of Hive 0.8.0) Assumes given timestamp is in given timezone and converts to UTC (as of Hive 0.8.0)

Hive

There is no direct way to check the version. Run the below command to check the version.

```
hive> set system:sun.java.command;
```

Move a file to hdfs first

```
ls ~/labs/map-reduce/hive-lab/free-txt-files/india.txt
hdfs dfs -mkdir labs
hdfs dfs -put ~/labs/map-reduce/hive-lab/free-txt-files/india.txt labs/
hdfs dfs -ls labs/india.txt
```

A simple hive code - word count example

- 1) connect to hive shell
- 2) create database labs;
- 3) create tables docs (line

```
CREATE TABLE docs (line STRING);
```

```
LOAD DATA INPATH 'docs' OVERWRITE INTO TABLE docs;
```

```
CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
(SELECT explode(split(line, '\s')) AS word FROM docs) w
GROUP BY word
ORDER BY word;
```

Hive is not a standalone tool and relies on various components for storing and querying data.

Within the Hadoop ecosystem, Hive is considered a client data access tool. Data access requires a compute, storage, management, as well as a security framework

Topic 1: Primitive Data Types

Type	Size	Example
.....
TINYINT	1 byte signed integer.	20
SMALLINT	2 byte signed integer.	20
INT	4 byte signed integer.	20
BIGINT	8 byte signed integer.	20
BOOLEAN	Boolean true or false.	TRUE
FLOAT	Single precision floating point.	3.14159
DOUBLE	Double precision floating point.	3.14159
STRING	Sequence of characters.	'Now is the time'
BINARY	Array of bytes. See discussion below	

It's useful to remember that each of these types is implemented in Java.

So the particular behavior details will be exactly what you would expect from the corresponding Java types.

For example, STRING is implemented by the Java String, FLOAT is implemented by Java float, etc.

Values of the new TIMESTAMP type can be integers, which are interpreted as seconds since the Unix epoch time (Midnight, January 1, 1970), floats, which are interpreted as seconds

since the epoch time with nanosecond resolution (up to 9 decimal places), and strings, which are interpreted according to the JDBC date string format convention,

YYYY-MM-DD hh:mm:ss.fffffffff.

Collection Data Types

STRUCT

Analogous to a C struct or an “object.”

`struct('John', 'Doe')`

Fields can be accessed using the “dot” notation.

For example, if a column name is of type STRUCT

`{first STRING; last STRING}`, then the

first name field can be referenced using `name.first`.

MAP

A collection of key-value tuples,

`map('first', 'John', 'last', 'Doe')`

where the fields are accessed using array notation (e.g., `['key']`).

For example, if a column name is of type MAP with key→value pairs

`'first'→'John'` and `'last'→'Doe'`, then the last name can be

referenced using `name['last']`.

ARRAY

Ordered sequences of the same type that are indexable using

`array('John', 'Doe')`

zero-based integers. For example, if a column name is of type

ARRAY of strings with the value `['John', 'Doe']`, then

the second element can be referenced using `name[1]`.

Most relational databases don’t support such collection types, because using them tends to break normal form.

For example, in traditional data models, structs might be captured in separate tables, with foreign key relations between the tables, as appropriate.

A practical problem with breaking normal form is the greater risk of data duplication, leading to unnecessary disk space consumption and potential data inconsistencies, as duplicate copies can grow out of sync as changes are made

Here is a table declaration that demonstrates how to use these types, an employees table in a fictitious Human Resources application:

```
CREATE TABLE employees (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT>);
```

Hive's default record and field delimiters

- 1) \n => For text files, each line is a record, so the line feed character separates records.
- 2) ^A ("control" A) Separates all fields (columns). Written using the octal code \001 when explicitly specified in CREATE TABLE statements.
- 2) ^B Separate the elements in an ARRAY or STRUCT, or the key-value pairs in a MAP. Written using the octal code \002 when explicitly specified in CREATE TABLE statements.
- 3) ^C Separate the key from the corresponding value in MAP key-value pairs. Written using the octal code \003 when explicitly specified in CREATE TABLE statements.

```
CREATE TABLE employees (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT>
)
```

ROW FORMAT DELIMITED

```
FIELDS TERMINATED BY '\001'
COLLECTION ITEMS TERMINATED BY '\002'
MAP KEYS TERMINATED BY '\003'
LINES TERMINATED BY '\n'
STORED AS TEXTFILE;
```

Note :

The ROW FORMAT DELIMITED sequence of keywords must appear before any of the other clauses.

The character \001 is the octal code for ^A. The clause ROW FORMAT DELIMITED FIELDS TERMINATED BY '\001' means that Hive will use the ^A character to separate fields.

Similarly, the character \002 is the octal code for ^B. The clause ROW FORMAT DELIMITED COLLECTION ITEMS TERMINATED BY '\002' means that Hive will use the ^B character to separate collection items.

Finally, the character \003 is the octal code for ^C. The clause ROW FORMAT DELIMITED MAP KEYS TERMINATED BY '\003' means that Hive will use the ^C character to separate map keys from values.

Database in Hive & Data Definition

- 1) HiveQL is the Hive query language.
- 2) Like all SQL dialects in widespread use, it doesn't fully conform to any particular revision of the ANSI SQL standard.
- 3) It is perhaps closest to MySQL's dialect.
- 4) Hive offers no support for rowlevel inserts, updates, and deletes. Hive doesn't support transactions.
- 5) If you don't specify a database, the default database is used.
- 6) The Hive concept of a database is essentially just a catalog or namespace of tables

Lets Practice

1) CREATE DATABASE sales;

2) CREATE DATABASE IF NOT EXISTS sales;

3) SHOW DATABASES;

4) SHOW DATABASES LIKE 's.%';

5) hive> CREATE DATABASE IF NOT EXISTS sales

> COMMENT 'Holds all financial tables';

6) DESCRIBE DATABASE sales;

7) you can associate key-value properties with the database, although their only function currently is to provide a way of adding information to the output of DESCRIBE DATABASE EXTENDED <database>:

hive> CREATE DATABASE sales7

> WITH DBPROPERTIES ('creator' = 'Mark Moneybags', 'date' = '2012-01-02');

8) DESCRIBE DATABASE EXTENDED sales7;

9) user sales - is the command to use the database, however there is no command to see which database is being used.

10) Setting a property to print the current database as part of the prompt

hive> set hive.cli.print.current.db=true;

11) DROP DATABASE IF EXISTS sales;

The IF EXISTS is optional and suppresses warnings if financials doesn't exist.

By default, Hive won't permit you to drop a database if it contains tables.

You can either drop the tables first or append the CASCADE keyword to the command, which will cause the Hive to drop the tables in the database first:

hive> DROP DATABASE IF EXISTS financials CASCADE;

When a database is dropped, its directory is also deleted

12) Alter Database

12.1) You can set key-value pairs in the DBPROPERTIES associated with a database using the ALTER DATABASE command.

```
hive> ALTER DATABASE sales SET DBPROPERTIES ('edited-by' = 'Labs');  
Imp **** There is no way to delete or "unset" a DBPROPERTY.
```

13) Creating Tables

```
CREATE TABLE IF NOT EXISTS sales.employees (  
    name STRING COMMENT 'Employee name',  
    salary FLOAT COMMENT 'Employee salary',  
    subordinates ARRAY<STRING> COMMENT 'Names of subordinates',  
    deductions MAP<STRING, FLOAT>  
        COMMENT 'Keys are deductions names, values are percentages',  
    address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT>  
        COMMENT 'Home address'  
    COMMENT 'Description of the table'  
    TBLPROPERTIES ('creator'='me', 'created_at'='2012-01-02 10:00:00', ...)  
    LOCATION '/user/hive/warehouse/mydb.db/employees';
```

Imp Note

* Hive automatically adds two table properties: last_modified_by holds the username of the last user to modify the table, and last_modified_time holds the epoch time in seconds of that modification.

* show tblproperties employees

14) You can also copy the schema (but not the data) of an existing table:

```
CREATE TABLE IF NOT EXISTS sales.employees2  
    LIKE sales.employees;
```

15) see the tables

```
show tables;  
show tables emp.*;  
show tables in sales;
```

16) Describe the table

```
DESCRIBE EXTENDED sales.employees;  
DESCRIBE FORMATTED sales.employees;
```

17) If you only want to see the schema for a particular column, append the column to the table name

```
describe sales.employee
```

Managed Tables in Hive

The tables we have created so far are called managed tables or sometimes called internal tables, because Hive controls the lifecycle of their data (more or less). As we've seen,

Hive stores the data for these tables in a subdirectory under the directory defined by `hive.metastore.warehouse.dir` (e.g., `/user/hive/warehouse`)

```
run this command and see hdfs dfs -ls /user/hive/warehouse
```

When we drop a managed table, Hive deletes the data in the table.

External Table

Suppose we are analyzing data from the stock markets. Periodically, we ingest the data for NASDAQ and the NYSE.

<https://www.nasdaq.com/symbol/csv/historical>

```
CREATE EXTERNAL TABLE IF NOT EXISTS stocks (
```

```
    date STRING,  
    close FLOAT,  
    volume INT,
```

```
open FLOAT,  
high FLOAT,  
low FLOAT)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','  
LOCATION '/landing/stocks';
```

** The EXTERNAL keyword tells Hive this table is external and the LOCATION ... clause is required to tell Hive where it's located.

** Because it's external, Hive does not assume it owns the data. Therefore, dropping the table does not delete the data, although the metadata for the table will be deleted.

Copy schema but change the data set

```
CREATE EXTERNAL TABLE IF NOT EXISTS sales.employees3  
LIKE sales.employees  
LOCATION '/path/to/data';
```

** If you omit the EXTERNAL keyword and the original table is external, the new table will also be external

** If you omit EXTERNAL and the original table is managed, the new table will also be managed.

Alter Tables Statements

Most table properties can be altered with ALTER TABLE statements, which change metadata about the table but not the data itself.

Renaming a Table : ALTER TABLE employee RENAME TO xemployee;

Changing Column :

```
ALTER TABLE employee  
CHANGE COLUMN salary total_salary INT  
COMMENT 'changing it'  
AFTER deduction;
```

Adding Column :

```
ALTER TABLE log_messages ADD COLUMNS (
    app_name STRING COMMENT 'Application name',
    session_id LONG COMMENT 'The current session id');
```

Deleting or Replacing Columns :

```
ALTER TABLE log_messages REPLACE COLUMNS (
    hours_mins_secs INT COMMENT 'hour, minute, seconds from timestamp',
    severity STRING COMMENT 'The message severity'
    message STRING COMMENT 'The rest of the message');

ALTER TABLE log_messages ADD IF NOT EXISTS
PARTITION (year = 2011, month = 1, day = 1) LOCATION '/logs/2011/01/01'
PARTITION (year = 2011, month = 1, day = 2) LOCATION '/logs/2011/01/02'
PARTITION (year = 2011, month = 1, day = 3) LOCATION '/logs/2011/01/03'
```

Alter Storage Properties

```
ALTER TABLE log_messages
PARTITION(year = 2012, month = 1, day = 1)
SET FILEFORMAT SEQUENCEFILE;
```

Inserting Data Into Hive Table

Partitioned, Managed Tables

It's used for distributing load horizontally, moving data physically closer to its most frequent users, and other purposes

```
CREATE TABLE employees (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT>
)
PARTITIONED BY (country STRING, state STRING);
```

Partitioning tables changes how Hive structures the data storage. If we create this table in the mydb database, there will still be an employees directory for the table:

```
...
.../employees/country=CA/state=AB
.../employees/country=CA/state=BC
...
.../employees/country=US/state=AL
.../employees/country=US/state=AK
...
```

You can see the partitions that exist with the SHOW PARTITIONS command:

```
hive> SHOW PARTITIONS employees;
```

Key Configuration

- 1) The database directory is created under a top-level directory specified by the property `hive.metastore.warehouse.dir`
- 2) You can override this default location for the new directory as shown in this example `CREATE DATABASE sales LOCATION '/my/preferred/directory';`

Complex Structure

```
-- define some default properties
```

```
set hive.cli.print.header=true;
```

```
set hive.cli.print.header=true;
```

```
-- Complex Data Type Example
```

```
drop table if exists employee_internal;
```

```
CREATE TABLE employee_internal
```

```
(
```

```
    name string,
```

```
    work_place ARRAY<string>,
```

```
    gender_age STRUCT<gender:string,age:smallint>,
```

```
    skills_score MAP<string,smallint>,
```

```
    depart_title MAP<string,ARRAY<string>>
```

```
)
```

```
COMMENT 'This is an internal table'
```

```
ROW FORMAT DELIMITED
```

```
  FIELDS TERMINATED BY '|'
```

```
  COLLECTION ITEMS TERMINATED BY ','
```

```
  MAP KEYS TERMINATED BY ':'
```

```
  STORED AS TEXTFILE;
```

```
describe formatted employee_internal;
```

Sample Data File : employee.txt

```
Michael|Montreal,Toronto|Male,30|DB:80|Product:Developer^DLead
```

```
Will|Montreal|Male,35|Perl:85|Product:Lead,Test:Lead
```

Shelley|New York|Female,27|Python:80|Test:Lead,COE:Architect

Lucy|Vancouver|Female,57|Sales:89,HR:94|Sales:Lead

```
LOAD DATA LOCAL INPATH 'labs/map-reduce/hive-lab/complex-data/employee.txt'
```

```
OVERWRITE INTO TABLE employee_internal;
```

```
LOAD DATA LOCAL INPATH 'employee3.txt' OVERWRITE INTO TABLE employee_internal;
```

--get all the data

```
SELECT * FROM employee_internal;
```

--get work place which is array of string

```
SELECT name, work_place FROM employee;
```

```
SELECT name, work_place[0] AS col_1, work_place[1] AS col_2, work_place[2] AS col_3 FROM employee_internal;
```

-- get the gender and age structure

```
SELECT name, gender_age FROM employee_internal;
```

-- get the gender and age as separate column

```
SELECT name, gender_age.gender, gender_age.age FROM employee_internal;
```

--get the skill score

```
SELECT skills_score FROM employee_internal;
```

```
SELECT name, skills_score['DB'] AS DB, skills_score['Perl'] AS Perl, skills_score['Python'] AS Python, skills_score['Sales'] as Sales, skills_score['HR'] as HR FROM employee_internal;
```

```
SELECT name, depart_title['Product'] AS Product, depart_title['Test'] AS Test, depart_title['COE'] AS COE, depart_title['Sales'] AS Sales FROM employee_internal;
```

-- **External Table**

```
drop table if exists employee_external;
CREATE TABLE employee_external
(
    name string,
    work_place ARRAY<string>,
    gender_age STRUCT<gender:string,age:smallint>,
    skills_score MAP<string,smallint>,
    depart_title MAP<string,ARRAY<string>>
)
COMMENT 'This is an internal table'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '|'
COLLECTION ITEMS TERMINATED BY ','
MAP KEYS TERMINATED BY ':'
STORED AS TEXTFILE
LOCATION '/user/cloudera/labs/map-reduce/hive-lab/complex-data/employee.txt';
```

-- Other supported format

-- SEQUENCEFILE, RCFILE, ORC, AVRO

-- **Partitioned**

```
drop table if exists employee_partitioned;
CREATE TABLE employee_partitioned (
    name string,
    work_place ARRAY<string>,
    gender_age STRUCT<gender:string,age:smallint>,
```

```
skills_score MAP<string,int>,
depart_title MAP<STRING,ARRAY<STRING>>
)
PARTITIONED BY (Year INT, Month INT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '|'
COLLECTION ITEMS TERMINATED BY ','
MAP KEYS TERMINATED BY ':';
```

```
SHOW PARTITIONS employee_partitioned;
```

```
-- this does not result any data
```

The partition is not enabled automatically. We have to use ALTER TABLE ADD PARTITION to add partitions to a table.

```
ALTER TABLE employee_partitioned ADD
PARTITION (year=2017, month=10)
PARTITION (year=2017, month=11);
```

```
SHOW PARTITIONS employee_partitioned;
```

-- Drop a partition

```
ALTER TABLE employee_partitioned
DROP IF EXISTS PARTITION (year=2017, month=10);
```

Sample Data File : employee_partitioned.txt

```
Michael|Montreal,Toronto|Male,30|DB:80|Product:Developer^DLead
Will|Montreal|Male,35|Perl:85|Product:Lead,Test:Lead
Shelley|New York|Female,27|Python:80|Test:Lead,COE:Architect
Lucy|Vancouver|Female,57|Sales:89,HR:94|Sales:Lead
```

```
LOAD DATA LOCAL INPATH 'employee_partitioned.txt' OVERWRITE INTO TABLE employee_partitioned PARTITION (year=2017, month=10);
```

```
SELECT name, year, month FROM employee_partitioned;
```

```
describe formatted employee_external;
```

BUCKET

Sample Data File : employee_id.txt

```
100|Michael|Montreal,Toronto|Male,30|DB:80|Product:Developer^DLead  
100|Will|Montreal|Male,35|Perl:85|Product:Lead,Test:Lead  
100|Shelley|New York|Female,27|Python:80|Test:Lead,COE:Architect  
100|Lucy|Vancouver|Female,57|Sales:89,HR:94|Sales:Lead  
100|Michael2|Montreal,Toronto|Male,30|DB:80|Product:Developer^DLead  
100|Will2|Montreal|Male,35|Perl:85|Product:Lead,Test:Lead  
200|Shelley2|New York|Female,27|Python:80|Test:Lead,COE:Architect  
200|Lucy2|Vancouver|Female,57|Sales:89,HR:94|Sales:Lead  
200|Michael3|Montreal,Toronto|Male,30|DB:80|Product:Developer^DLead  
200|Will3|Montreal|Male,35|Perl:85|Product:Lead,Test:Lead  
200|Shelley3|New York|Female,27|Python:80|Test:Lead,COE:Architect  
200|Lucy3|Vancouver|Female,57|Sales:89,HR:94|Sales:Lead
```

```
drop table if exists employee_id;
```

```
CREATE TABLE employee_id
```

```
(
```

```
emp_id int,
```

```
name string,
```

```
work_place ARRAY<string>,
```

```
gender_age STRUCT<gender:string,age:smallint>,
skills_score MAP<string,smallint>,
depart_title MAP<string,ARRAY<string>>
)
COMMENT 'This is an internal table'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '|'
COLLECTION ITEMS TERMINATED BY ','
MAP KEYS TERMINATED BY ':'
STORED AS TEXTFILE;
```

```
LOAD DATA LOCAL INPATH 'employee_id.txt' OVERWRITE INTO TABLE employee_id ;
```

```
drop table if exists employee_id_bucket;
CREATE TABLE employee_id_bucket
(
emp_id int,
name string,
work_place ARRAY<string>,
gender_age STRUCT<gender:string,age:smallint>,
skills_score MAP<string,smallint>,
depart_title MAP<string,ARRAY<string>>
)
COMMENT 'This is an internal table'
CLUSTERED BY (emp_id) INTO 2 BUCKETS
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '|'
COLLECTION ITEMS TERMINATED BY ','
MAP KEYS TERMINATED BY ':'
STORED AS TEXTFILE;
```

```
-- Note: if clustered column string does not match with any of the column, it may throw error;  
set map.reduce.tasks = 2;  
set hive.enforce.bucketing = true;
```

```
INSERT OVERWRITE TABLE employee_id_bucket SELECT * FROM employee_id;
```

View

```
CREATE VIEW employee_skills  
AS  
SELECT name, skills_score['DB'] AS DB,  
skills_score['Perl'] AS Perl,  
skills_score['Python'] AS Python,  
skills_score['Sales'] as Sales,  
FROM employee;  
No rows affected (0.253 seconds)
```

Json Example

Schema of the data for JSON is

- 1) blogID
- 2) date
- 3) commenter name
- 4) comment
- 5) commenter email
- 6) commenter web site

Sample data comment.txt

```
{ "blogID" : "FJY26J1333", "date" : "2012-04-01", "name" : "vpxnksu", "comment" : "good stuff",
"contact" : { "email" : "vpxnksu@gmail.com", "website" : "vpxnksu.wordpress.com" } }

{ "blogID" : "FJY26J1333", "date" : "2012-04-01", "name" : "vpxnksu", "comment" : "good stuff",
"contact" : { "email" : "vpxnksu@gmail.com", "website" : "vpxnksu.wordpress.com" } }

{ "blogID" : "VSAUMDFGSD", "date" : "2012-04-01", "name" : "yhftrcx", "comment" : "another
comment",}
```

```
CREATE EXTERNAL TABLE comments_external
(
  cmt STRING
)
COMMENT 'This is an external table'
ROW FORMAT DELIMITED
STORED AS TEXTFILE
LOCATION '/user/test/comment';
```

```
LOAD DATA LOCAL INPATH 'comment.txt' OVERWRITE INTO TABLE comments_external;
```

```
SELECT b.blogID, c.email FROM comments_external a LATERAL VIEW json_tuple('blogID', 'contact') b
AS blogID, contact LATERAL VIEW json_tuple(b.contact, 'email', 'website') c
AS email, website WHERE b.blogID='64FY4D0B28';
```

PySpark

Spark Submit Example

```
#import necessary functions at the begining of the program
from pyspark.sql import SparkSession
from pyspark.sql import functions as F

#standard varialbes
APP_NAME = "My-Simple-PySpark-Program"

def main(spark_session):
    # Have you processing here
    # Create a datafram using range function which start with 1 and ends with 100 with 2 as step (1,3,
5 ...)
    df = spark_session.range(1,100,2)
    df.printSchema()
    df.show()

if __name__ == '__main__':
    print("The program started... ")
    #Create The spark session by giving app name

    #https://spark.apache.org/docs/latest/api/python/pyspark.sql.html#pyspark.sql.SparkSession
    spark = SparkSession.builder.appName(APP_NAME).getOrCreate()

    main(spark)
    print("Program is over")
```

Reading Json File

```
spark.read.json("file:///tmp/tips/data/flight-data.json")

...
DataFrame[DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string, count: bigint]
...

tips_df = spark.read.json("file:///tmp/tips/data/flight-data.json")
tips_df.printSchema()

...
root
 |-- DEST_COUNTRY_NAME: string (nullable = true)
 |-- ORIGIN_COUNTRY_NAME: string (nullable = true)
 |-- count: long (nullable = true)
...

tips_df.show()

...
+-----+-----+---+
| DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|
+-----+-----+---+
| United States| Romania| 15|
| United States| Croatia| 1|
| United States| Ireland| 344|
| Egypt| United States| 15|
| United States| India| 62|
| United States| Singapore| 1|
| United States| Grenada| 62|
| Costa Rica| United States| 588|
| Senegal| United States| 40|
| Moldova| United States| 1|
| United States| Sint Maarten| 325|
| United States| Marshall Islands| 39|
| Guyana| United States| 64|
| Malta| United States| 1|
| Anguilla| United States| 41|
| Bolivia| United States| 30|
| United States| Paraguay| 6|
| Algeria| United States| 4|
| Turks and Caicos ...| United States| 230|
```

```

|    United States|      Gibraltar|   1|
+-----+-----+----+
only showing top 20 rows
"""

tips_df =
spark.read.format("json").option("mode","FAILFAST").option("inferSchema","true").load("file:///
tmp/tips/data/flight-data.json")
tips_df.printSchema()

"""

root
|-- DEST_COUNTRY_NAME: string (nullable = true)
|-- ORIGIN_COUNTRY_NAME: string (nullable = true)
|-- count: long (nullable = true)
"""

tips_df.createOrReplaceTempView("tips_table")
tips_df.createOrReplaceTempView("tips_table")
spark.sql("select count from tips_table").show(5)

"""

20/05/07 10:19:55 WARN ObjectStore: Version information not found in metastore.
hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0
20/05/07 10:19:55 WARN ObjectStore: Failed to get database default, returning
NoSuchObjectException
20/05/07 10:19:56 WARN ObjectStore: Failed to get database global_temp, returning
NoSuchObjectException
+----+
|count|
+----+
|  15|
|   1|
| 344|
|  15|
|  62|
+----+
only showing top 5 rows
"""

```

Rearrange Column In Dataframe

```
#import necessary files
from pyspark.sql.types import StructType, StructField, IntegerType, StringType

tips_list = [[1,"item 1" , 2], [2 , "item 2", 3], [1 , "item 3" ,2], [3 , "item 4" , 5]]
tips_schema = StructType( [StructField("Item_ID",IntegerType(),True),StructField("Item_Name",StringType(), True ), StructField("Quantity",IntegerType(), True)]) 

#create dataframe using spark session
tips_df = spark.createDataFrame(tips_list, tips_schema)

#the schema
tips_df.printSchema()

"""

root
 |-- Item_ID: integer (nullable = true)
 |-- Item_Name: string (nullable = true)
 |-- Quantity: integer (nullable = true)
"""

#register as table
tips_df.createOrReplaceTempView("tips_table")

tips_new_df = spark.sql("select Item_Name, Item_ID, Quantity from tips_table")

tips_new_df.printSchema()

"""

root
 |-- Item_Name: string (nullable = true)
 |-- Item_ID: integer (nullable = true)
 |-- Quantity: integer (nullable = true)
"""

#using the dataframe approach

#create dataframe using spark session
tips_df = spark.createDataFrame(tips_list, tips_schema)

#use the select function from the dataframe
tips_new_df = tips_df.select("Item_Name","Item_ID","Quantity")
tips_new_df.printSchema()
```

